

# Enzo-P/Cello


## Computational Cosmology and Astrophysics on Adaptive Meshes using Charm++

James Bordner, Michael L. Norman

University of California, San Diego  
San Diego Supercomputer Center

Supercomputing 2018  
2018-11-11/16

# Scientific questions in astrophysics and cosmology



[ John Wise ]

# The ENZO MPI-parallel AMR astrophysics application

## ■ ENZO is a very powerful application


- science: applicable to wide range of problems
- equations: rich multiphysics capabilities
- methods: wide range of numerical solvers
- data structures: SAMR, fields + particles

## ■ Its main limitation is parallel scalability

- ENZO was developed starting in 1994
- “massive parallelism” meant  $P \approx 10^3$
- $P \approx 10^7$  today

## ■ Motivated AMR data structure redesign

- Enzo-P: “Petascale” branch of ENZO
- keep ENZO’s physics and many methods
- Cello highly scalable AMR framework



[ Sam Skillman, Matt Turk ]

# The ENZO MPI-parallel AMR astrophysics application

## ■ ENZO is a very powerful application


- science: applicable to wide range of problems
- equations: rich multiphysics capabilities
- methods: wide range of numerical solvers
- data structures: SAMR, fields + particles

## ■ Its main limitation is parallel scalability

- ENZO was developed starting in 1994
- “massive parallelism” meant  $P \approx 10^3$
- $P \approx 10^7$  today

## ■ Motivated AMR data structure redesign

- Enzo-P: “Petascale” branch of ENZO
- keep ENZO’s physics and many methods
- Cello highly scalable AMR framework



[ Sam Skillman, Matt Turk ]

# The ENZO MPI-parallel AMR astrophysics application

## ■ ENZO is a very powerful application


- science: applicable to wide range of problems
- equations: rich multiphysics capabilities
- methods: wide range of numerical solvers
- data structures: SAMR, fields + particles

## ■ Its main limitation is parallel scalability

- ENZO was developed starting in 1994
- “massive parallelism” meant  $P \approx 10^3$
- $P \approx 10^7$  today

## ■ Motivated AMR data structure redesign

- **Enzo-P**: “Petascale” branch of ENZO
- keep ENZO’s physics and many methods
- **Cello** highly scalable AMR framework





[ Sam Skillman, Matt Turk ]

# Cello scalable adaptive mesh refinement

## Key differences between ENZO and Enzo-P

### Enzo-P/Cello

- array of octrees AMR
- Charm++ parallelization
- reusable AMR framework



### ENZO



- structured AMR
- MPI+OpenMP parallelization
- non-encapsulated AMR data structure

# Cello scalable adaptive mesh refinement

## Key differences between ENZO and Enzo-P

### Enzo-P/Cello

- array of octrees AMR
- Charm++ parallelization
- reusable AMR framework



### ENZO



- structured AMR
- MPI+OpenMP parallelization
- non-encapsulated AMR data structure

# Cello scalable adaptive mesh refinement

## Key differences between ENZO and Enzo-P

### Enzo-P/Cello

- array of octrees AMR
- Charm++ parallelization
- reusable AMR framework



### ENZO

- structured AMR
- MPI+OpenMP parallelization
- non-encapsulated AMR data structure



# Cello distributed adaptive mesh refinement


How field data are communicated between blocks

## ■ Data-driven execution

- send Field face data when available
- indexed using bit-coding in hierarchy
- count face data received
- last receive triggers computation

## ■ Dynamic task scheduling



- multiple Blocks per process
- overlapped communication/computation



# Enzo-P/Cello NSF Blue Waters scaling

“Alphabet Soup” test: hydro and particles



## We tested basic Enzo-P hydrodynamics and particles scalability



- variation of “array of Sedov Blast” test
- letters instead of spheres
  - inhibits lockstep coarsen/refine
- one blast per BW fp-core
- tested with/without tracer particles
- $32^3$  or  $24^3$  cells per block
- decent sized AMR problem for 2016
  - 262K fp-cores
  - 50M Blocks
  - 1.7T cells; 0.7T (cells + particles)
- ENZO would need 72GB per process!

# Enzo-P/Cello NSF Blue Waters scaling


“Alphabet Soup” test: hydro and particles



# Enzo-P/Cello NSF Blue Waters scaling

“Unigrid Cosmology” test: hydro, particles, gravity

## We tested scaling of more recent support for cosmology





[ Renyue Cen ]

- PPM hydrodynamics
- “dark matter” particles
- CIC particle-mesh gravity
- multigrid solver—“unigrid” only
- tested up to 131K fp-cores
- have since implemented AMR solver
  - Dan Reynolds “HG” algorithm
  - MG preconditioned BiCG-STAB
  - “semi-scalable”
  - DD and AFACx MG solvers soon

# Enzo-P/Cello NSF Blue Waters scaling

“Unigrid Cosmology” test: hydro, particles, gravity



# Conclusions and next steps

## Enzo-P/Cello is a highly scalable branch of ENZO

- Charm++ enables fully-distributed AMR mesh hierarchy
- excellent hydro weak scaling through  $P = 262K$
- very good “unigrid” cosmology scaling through  $P = 131K$
- scalable AMR cosmology soon ( $< 1$  month)

## Next steps include

- improve I/O scalability
- implement block-adaptive time-stepping
- run large-scale AMR cosmology simulations
- next phase: *Enzo-E*
  - improve strong scaling
  - support heterogeneous hardware
  - increase ENZO developer engagement

# Acknowledgements

Funding for development of Enzo-P/Cello has been provided by the National Science Foundation grants OAC-1835402, SI2-SSE-1440709, PHY-1104819, and AST-0808184.


This research is part of the Blue Waters sustained-petascale computing project, which is supported by the National Science Foundation (awards OCI-0725070 and ACI-1238993) and the state of Illinois. Blue Waters is a joint effort of the University of Illinois at Urbana-Champaign and the National Center for Supercomputing Applications.

<http://cello-project.org/>

# Cello distributed adaptive mesh refinement

How particle data are communicated between Blocks


- Communication is required when particles move outside a Block
- This is done using a 4x4x4 array
  - array contains pointers to ParticleData (PD) objects
  - one PD object per neighbor Block



- Migrating particles are
  - scatter()-ed to PD array objects
  - sent to associated neighbors
  - gather()-ed by neighbors
- One sweep through particles
- One communication step per neighbor




# The Charm++ parallel programming system



## A Charm++ parallel program

- Charm++ program
  - Decomposed by *objects*
  - Charm++ objects called *chares*
  - invoke *entry methods*
  - *asynchronous*
  - communicate via *messages*
- Charm++ runtime system
  - maps chares to processors
  - schedules entry methods
  - can migrate chares
- Additional features
  - checkpoint/restart
  - dynamic load balancing
  - fault-tolerance


# The Charm++ parallel programming system



## A Charm++ parallel program

- Charm++ program
  - Decomposed by *objects*
  - Charm++ objects called *chares*
  - invoke *entry methods*
  - *asynchronous*
  - communicate via *messages*
- Charm++ runtime system
  - maps chares to processors
  - schedules entry methods
  - can migrate chares
- Additional features
  - checkpoint/restart
  - dynamic load balancing
  - fault-tolerance

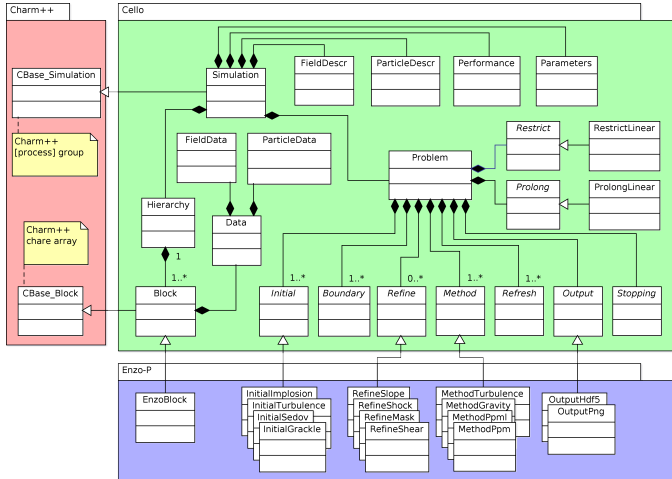
# The Charm++ programming system




## A Charm++ parallel program

- Charm++ program
  - Decomposed by *objects*
  - Charm++ objects called *chares*
  - invoke *entry methods*
  - *asynchronous*
  - communicate via *messages*
- Charm++ runtime system
  - maps chares to processors
  - schedules entry methods
  - can migrate chares
- Additional features
  - checkpoint/restart
  - dynamic load balancing
  - fault-tolerance


# Enzo-P/Cello/Charm++ class organization




# How the Block chare array is indexed



- User-defined chare array indices supported
- Cello indices for Block arrays:
  - $3 \times 10$  bits for *array indices*
  - $3 \times 20$  bits for the *octree encoding*
  - 6 bits for the *block level*
- Up to  $1024^3$  array of octrees
- Up to 20 octree levels
- $-31 \leq \text{level} \leq 31$
- Block id's use index: e.g. `B100:11_1:01`



# How Particle objects store particle data



- multiple particle *types*
- particles allocated in *batches*
  - fixed size arrays
  - fewer new/delete operations
  - efficient insert/delete operations
  - potentially useful for GPU's
- batches store particle *attributes*
  - (position, velocity, mass, etc.)
  - 8,16,32,64-bit integers
  - 32,64,128-bit floats

- particle positions may be floating-point or integers
  - floating-point for storing global positions
  - integers for Block-local coordinates
    - solves reduced precision issue for deep hierarchies
    - less memory required for given accuracy

# How do Enzo-P and ENZO differ?

	ENZO	Enzo-P/Cello
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

# How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint



# How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

# How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

# How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

\* not implemented yet

# How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

\* not implemented yet

# How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

\* not implemented yet

## How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

\* not implemented yet

## How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

\* not implemented yet

## How do Enzo-P and ENZO differ?

	<b>ENZO</b>	<b>Enzo-P/Cello</b>
Parallelization	MPI	Charm++
AMR type	patch-based	octree-based
AMR structure	replicated	fully distributed
Time stepping	level-adaptive	block-adaptive*
Block sizes	×1000 variation	constant
Task scheduling	level-parallel	dependency-driven
Load balancing	patch migration	Charm++-based
Data locality	LB conflict	no LB conflict
Mesh quality	level jumps	2-to-1 constraint

\* not implemented yet