

Simulating Ultralight Dark Matter with Chapel

An Experience Report

Nikhil Padmanabhan ¹ Elliot Ronaghan ²

J. Luna Zagorac ¹ Richard Easter ³

2019/11/17

¹Yale Univ.

²Cray

³Univ. of Auckland

Table of contents

1. Introduction
2. FFTs in Chapel
3. Isolated Gravitational Potentials
4. The Full Simulation
5. Lessons Learned

Introduction

Motivating Ultralight Dark Matter

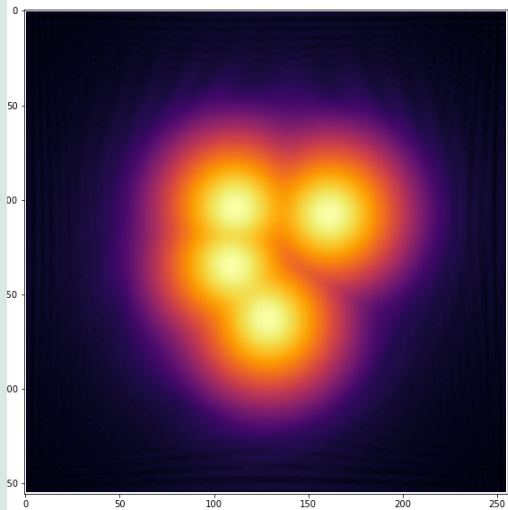
- In the standard cosmological model, 80% of the matter in the Universe is “dark” (i.e. non-baryonic).
- Explains a large scale of observations, from the rotation of galaxies, to “Bullet” clusters, to the distribution of galaxies, to the cosmic microwave background.
- Form gravitationally bound structures : dark matter halos.
- The traditional model is a heavy particle ($\sim 100\times$ proton), with weak interactions.
- Possible puzzles remain on small scales from the structure of dark matter halos, to the observed abundance of dark matter halos. Note that these might well be solved by astrophysics.
- We have not detected these in the lab, or at accelerators.

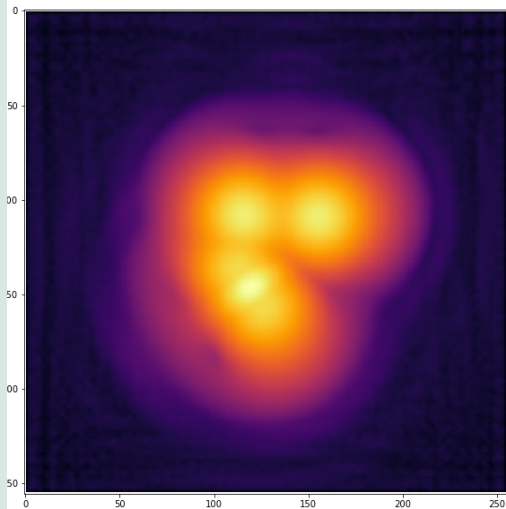
Motivating Ultralight Dark Matter

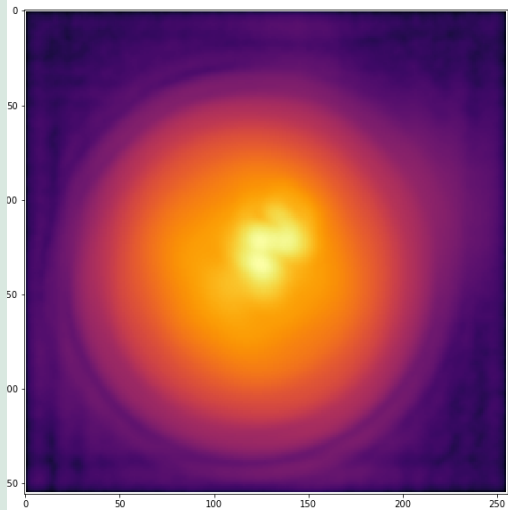
- A different paradigm is a very light particle ($10^{-31} \times$ proton).
- Many names : fuzzy dark matter, Bose-Einstein dark matter, ...
- Small mass means that quantum-mechanics can smear it out over astrophysically interesting scales.
- High enough density that it forms a Bose-Einstein condensate.
- Different phenomenology : eg. interference patterns.
- Anything by the most idealized situations requires simulations.

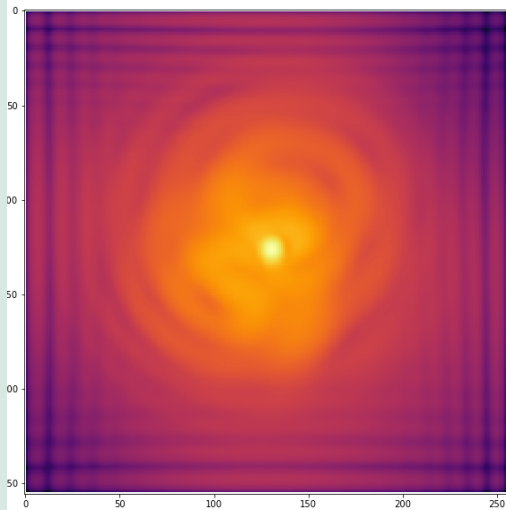
Our Motivation

- Want a code to do numerical experiments with.
- Need scalability
 - Must resolve soliton cores : large dynamic range
 - Simulation time scales as N^5 ; need to scale to large numbers of nodes.
- Initial problem : revisit aspects of the formation of ultra-light dark matter halos from collisions of soliton cores.
- This is an area of very active research (we are newcomers).
- Several codes exist - including adaptive codes, codes built on existing large astrophysical simulations. Challenges to large boxes still exist.
- An incomplete list : Schive et al, 2014, Mocz et al, 2017, Edwards et al 2017, Veltmaat et al, 2018









History of Project

- **PyUltraLight^a**: An initial code in Python, driven by Jupyter notebook
 - Easy to use and modify, allowing numerical experiments
 - Performant and multithreaded (made significant use of eg. **numexpr**, **FFTW**)
- Extending to isolated potentials hit Python bottlenecks
- Attempted a skunkworks (2019/6/22) port to Chapel for a single node. Resulting code not much longer than Python, could implement isolated potentials, better multithreaded performance.
- Distributed Code
 - Want to run larger N_{grid} , can we extend the code?
 - Isolated potential calculation led to wanting a native Chapel distributed FFT (useful for many other tasks).^b
 - Validating the FFT led to the NAS NPB benchmark.

^aEdwards et al, arXiv:1807.04037

^bNote that Chapel can also interoperate with MPI.

The Schrodinger-Poisson Equations

The Model

$$i\hbar \frac{\partial \psi}{\partial t} = -\frac{\hbar^2}{2m} \nabla^2 \psi + m\Phi \psi$$
$$\nabla^2 \Phi = 4\pi G m |\psi|^2$$

Timestepping

$$\psi(\mathbf{x}, t+h) \approx \exp\left[\frac{-ih}{2}\Phi(\mathbf{x}, t+h)\right] \exp\left[\frac{ih}{2}\nabla^2\right] \exp\left[\frac{-ih}{2}\Phi(\mathbf{x}, t)\right] \psi(\mathbf{x}, t)$$

Taking Fourier transforms simplifies the kinetic term

$$\psi(\mathbf{x}, t+h) \approx \exp\left[\frac{-ih}{2}\Phi(\mathbf{x}, t+h)\right] \mathcal{FT}^{-1} \exp\left[\frac{ih}{2}\mathbf{k}^2\right] \mathcal{FT} \exp\left[\frac{-ih}{2}\Phi(\mathbf{x}, t)\right] \psi(\mathbf{x}, t)$$

See eg. Edwards et al, arXiv:1807.04037

FFTs in Chapel

Slab Decompositions Are Simple

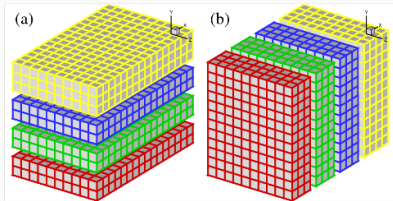


Figure 1: Slab Decomposition

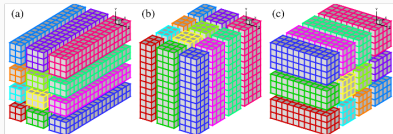


Figure 2: Pencil Decomposition

- Slab decompositions are simpler (especially for the end user)
- Slab limits the amount of parallelism expressed (especially with pure MPI)
- Use 1 slab per locale/node.
- Limits $N_{\text{grid}} \geq N_{\text{nodes}}$, but in practice, not limiting.
- Reduce communication complexity

<http://www.2decomp.org/decomp.html>

The Algorithm

1. Decompose array into slabs in the x direction
2. Fourier transform in the y direction¹
3. Fourier transform in the z direction
4. Transpose x and y (all to all)
5. Fourier transform in the x direction

¹We use FFTW (www.fftw.org) for 1D serial transforms.

Chapel Code is Expressive : A Naive Implementation

```
1  coforall loc in Locales do on loc {
2      ...
3      for ix in xSrc {
4          myplane = Src[{ix..ix, ySrc, zSrc}];
5          // Y-transform
6          forall iz in zSrc {
7              yPlan.execute(myplane[0, ySrc.first, iz]);
8          }
9          // Z-transform, offset to reduce comm congestion/collision
10         forall iy in offset(ySrc) {
11             zPlan.execute(myplane[0, iy, zSrc.first]);
12             // Transpose data into Dst
13             Dst[{iy..iy, ix..ix, zSrc}] = myplane[{0..0, iy..iy, zSrc}];
14         }
15     }
16     // Wait until all communication is complete
17     allLocalesBarrier.barrier();
18     // X-transform, similar to Y-transform
19     ...
20 }
```


Chapel Code is Expressive : A Naive Implementation

```
1  coforall loc in Locales do on loc {
2      ...
3      for ix in xSrc {
4          myplane = Src[{ix..ix, ySrc, zSrc}];
5          // Y-transform
6          forall iz in zSrc {
7              yPlan.execute(myplane[0, ySrc.first, iz]);
8          }
9          // Z-transform, offset to reduce comm congestion/collision
10         forall iy in offset(ySrc) {
11             zPlan.execute(myplane[0, iy, zSrc.first]);
12             // Transpose data into Dst
13             Dst[{iy..iy, ix..ix, zSrc}] = myplane[{0..0, iy..iy, zSrc}];
14         }
15     }
16     // Wait until all communication is complete
17     allLocalesBarrier.barrier();
18     // X-transform, similar to Y-transform
19     ...
20 }
```

- Can use SPMD programming when needed
- Data parallel features
- PGAS - simple communication
- User-defined iterators (eg. offset)

Chapel Code is Expressive : A Performant Implementation

```
1      ...
2      forall iy in offset(ySrc) {
3          zPlan.execute(myplane[0, iy, zSrc.first]);
4          // Transpose data into Dst, and copy the next Src slice into myplane
5          copy(Dst[...], myplane[...], myLineSize);
6          if (ix != xSrc.last) {
7              copy(myplane[...], Src[...], myLineSize);
8          }
9      }
10     ...
```

- Overlap computation and communication, even with blocking comm
- Use low-level communication primitives

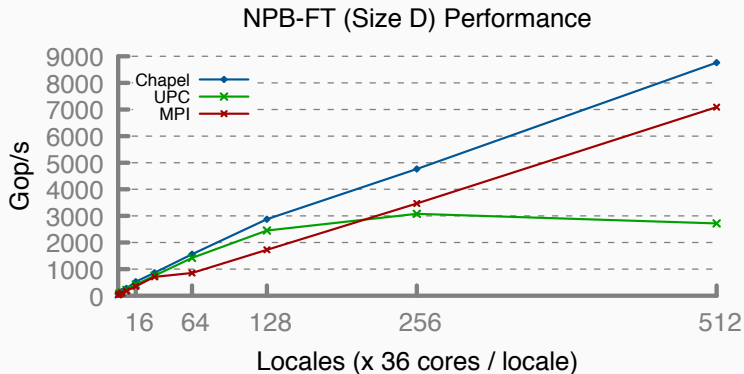
Machine/Compiler Specifications

- Scalability Hardware (Cray-XC):
 - 36-core (72 HT), 128 GB RAM
 - dual 18-core (36 HT) "Broadwell" 2.1 GHz processors
- Software
 - CLE 7.0.UP01
 - Intel Compilers 19.0.5.281
 - FFTW 3.3.8.4
 - Chapel 1.20.0
 - Cray 9.0.2 (classic)

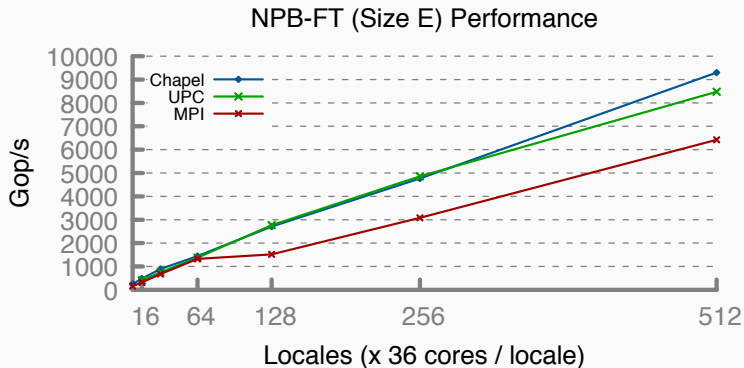
Benchmarks

- Use the NAS NPB-FT benchmark
 - NPB v3.4
 - Class D ($2048 \times 1024 \times 1024$), E ($8 \times$), F($8 \times$)
- Compare Chapel, MPI reference and UPC (with non-blocking overlapped comm)
- MPI and UPC use pencil decompositions for large problems/node counts.
- MPI and UPC use 32 cores/node (require a power of 2)
 - Restricting Chapel to 32 cores does not significantly change timings, indicating memory/communication bound.

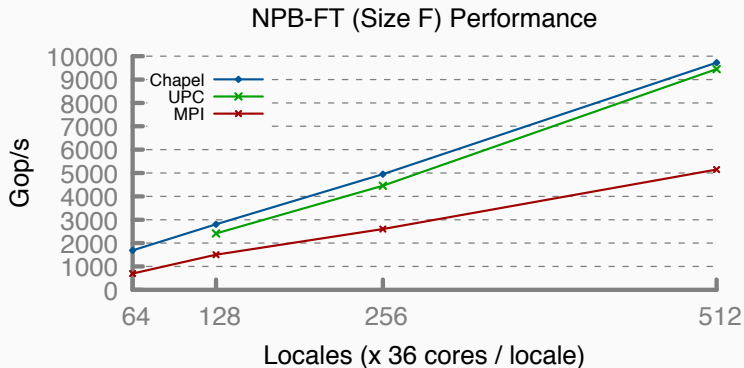
Chapel FFTs Scale Well Across Nodes : Class D



Chapel FFTs Scale Well Across Nodes : Class E = $8 \times D$

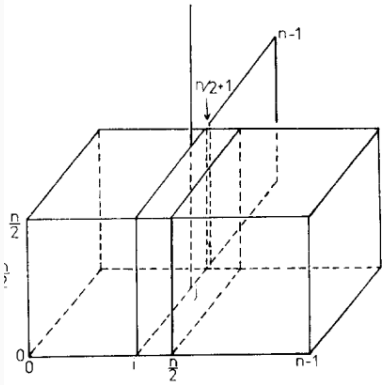


Chapel FFTs Scale Well Across Nodes : Class F = $8 \times E$



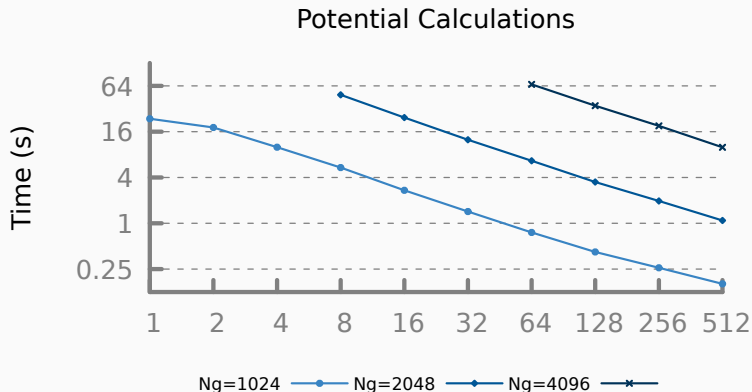
Isolated Gravitational Potentials

Calculating Potentials with Isolated Boundary Conditions



- Standard solution uses Fourier transforms, but assumes periodic boundary conditions.
- Need isolated boundary conditions $\Phi(\infty) = 0$.
- Convolve with $1/r$ Green's function, pad domain with zeros.
- Use separability of Fourier transforms to do each dimension separately.
- Only doubles the computational domain, not a factor of 8.

Potential Calculations Scale Well



The Full Simulation

Going to Distributed Code

Setting the boundary of the domain

```
1  iter boundary(d : domain, thick : int=1) {  
2      param rank = d.rank;  
3      for param idim in 1..rank {  
4          var off : rank*int;  
5          off(idim)=thick;  
6          yield d.interior(off);  
7  
8          off(idim)=-thick;  
9          yield d.interior(off);  
10     }  
11 }
```

- Same code as single locale version
- Sometimes, things work seamlessly!
- Highlights strength of Chapel domains

Going to Distributed Code

Initialization Code : Single-Node Version

```
1  proc addSoliton(..., psi : [?Dom] complex) {  
2      ...  
3      forall ijk in Dom {  
4          ...  
5          psi[ijk] += val*phase1*phase2;  
6      }  
7  }
```

Initialization Code : Distributed Version

```
1  proc addSoliton(..., psi : [?Dom] complex) {  
2      ...  
3      forall ijk in Dom with (var myprofile = this.profile,  
4                               ...)  
5      {  
6          ...  
7          psi.localAccess[ijk] += val*phase1*phase2;  
8      }
```

Going to Distributed Code

Initialization Code : Single-Node Version

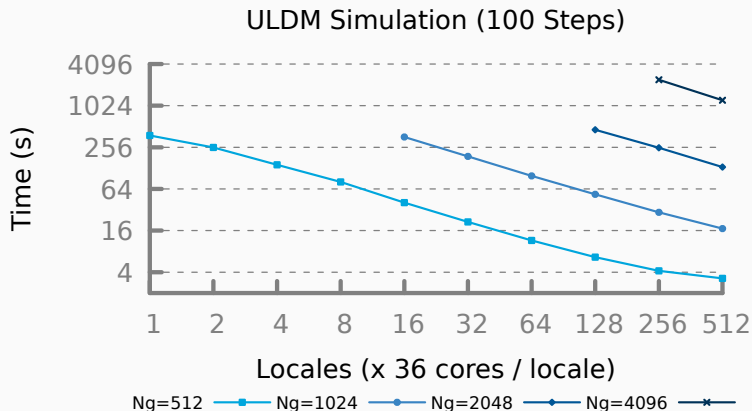
```
1  proc addSoliton(..., psi : [?Dom] complex) {  
2    ...  
3    forall ijk in Dom {  
4      ...  
5      psi[ijk] += val*phase1*phase2;  
6    }  
7  }
```

Initialization Code : Distributed Version

```
1  proc addSoliton(..., psi : [?Dom] complex) {  
2    ...  
3    forall ijk in Dom with (var myprofile = this.profile,  
4    ...) {  
5      {  
6        ...  
7        psi.localAccess[ijk] += val*phase1*phase2;  
8      }  
9    }
```

- Need to localize some variables (sometime will be surprised by comm)
- The compiler currently does not fully optimize local array accesses.
- Note that **psi** changes from a local to a distributed array, but this does not affect eg. the **forall** loop.

Simulations Scale Well



Lessons Learned

Where Chapel could do better

1. Tooling

- Identifying communication - how much and from where? How to recognize a sub-optimal pattern.
- Easier profiling
- Compiler improvements, including speed.

2. Easier to express low-level communication/locality

- Low level communication primitives are not exposed to user (useful when the user can reason better about the communication patterns).
- Verbose to express locality of computation and have the compiler optimize appropriately.

3. Fewer hidden performance traps

- Unexpected communication
- Promotion of operations over N-d arrays can be slow.

None of these are new issues to the Chapel team (and many have Github issues).

- HPC:
 - *Productivity*: Chapel design has scientific codes in mind.
 - Domains/Arrays
 - Expressive Parallelism - where/when you need it.
 - Interoperability - C (and now Python!)
 - *Performance*: Chapel code can perform/scale very well without heroic efforts.
- It's a fun language to write. Easy to throw together prototype code in. And it largely does the right thing!
- I'm getting to the point where I could imagine just working in Chapel. And I'm more open to getting students to work in it.
- This is the first code that I've had run on ~18K cores. Significant credit for that goes to Elliot (and the Chapel team), but the language played a non-trivial role here.

Backup Slides

Chapel FFTs Scale Well Across Nodes : Class D

