

Efficient Active Message RMA in GASNet Using a Target-Side Reassembly Protocol

Paul H. Hargrove and Dan Bonachea

`gasnet.lbl.gov`

`doi.org/10.25344/S4PC7M`

Parallel Applications Workshop, Alternatives to MPI+X (PAW-ATM)

Held in conjunction with SC19

November 17, 2019

Overview

- Introduce GASNet-EX
- Introduce Active Messages
- Discuss protocols for implementing Active Messages
- Measure a new “target-side reassembly” protocol for Active Messages on the Cray XC’s Aries network

GASNet-1: Overview

- Started in 2002 to provide a portable network communication runtime for three PGAS languages:
 - UPC, CAF and Titanium
- Primary features:
 - Non-blocking RMA (one-sided Put and Get)
 - Active Messages (simplification of Berkeley AM-2)
- Motivated by semantic issues in (then current) MPI-2.0
 - Dan Bonachea, Jason Duell, "Problems with using MPI 1.1 and 2.0 as compilation targets for parallel language implementations", IJHPCN 2004. doi.org/10.25344/S4JP4B

GASNet: Adoption and Portability

- Client runtimes

UPC++
Berkeley UPC
GCC/UPC
Clang UPC
Chapel

Legion
Titanium
Rice Co-Array Fortran
OpenUH Co-Array Fortran
OpenCoarrays in GCC Fortran

OpenSHMEM reference impl.
Omni XcalableMP
At least 7 others known to us

- Network conduits

OpenFabrics Verbs (InfiniBand)
Mellanox MXM and VAPI (InfiniBand)
Cray uGNI (Gemini and Aries)
Intel PSM2 (OmniPath)

IBM PAMI (BG/Q and others)
IBM DCMF (BG/P)
IBM LAPI (Colony and Federation)
Cray Portals3 (Seastar)

SHMEM (Cray X1 and SGI Altix)
Quadric elan3/4 (QsNet I/II)
Myricom GM (Myrinet)
Dolphin SISC

UDP (any TCP/IP network)
MPI 1.1 or newer

OFI/libfabric
Sandia Portals4

Shared memory (no network)

- Supported platforms

- Over 10 compiler families, 15 operating systems and dozens of architectures

* These lists and counts include both current and past support

GASNet-EX: Overview

- GASNet-EX is the next generation of GASNet
 - Addressing needs of newer programming models such as UPC++, Legion and Chapel
 - Incorporating over 15 years of lessons learned
 - Provides backward compatibility for GASNet-1 clients
- Motivating goals include
 - Support more client asynchrony
 - Enable more client adaptation
 - Decrease memory footprint
 - Improve threading support
 - Support offload to network h/w
 - Support multi-client applications
 - Support for device memory

GASNet-EX: Status

- GASNet-EX is still evolving
 - Not every new feature has been implemented yet
 - Most have - with benefits shown in prior work
- Four prominent clients actively adopting GASNet-EX
 - UPC++ and Berkeley UPC Runtime require GASNet-EX
 - Chapel embeds GASNet-EX
 - Legion has started work to use EX-specific features

GASNet-EX: New Features Include

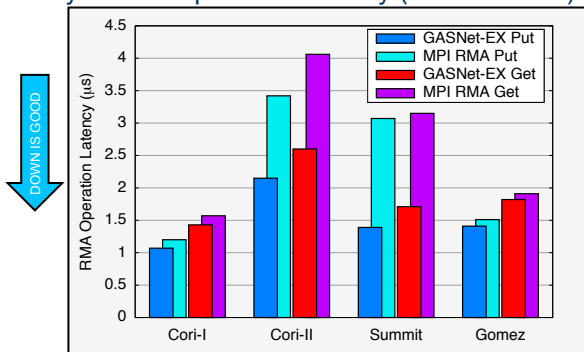
- Local completion control
 - Improved control over buffer lifetime to increase overlap
- Immediate-mode injection
 - Avoid stalls in low-resource conditions
- Negotiated-payload Active Messages
 - Construct messages in GASNet's buffers to avoid `memcpy()`
- Remote atomic operations
 - Utilize offload capabilities in modern network interfaces
- Subset teams
- Numerous small API additions and improvements

For details see Languages and Compilers for Parallel Computing (LCPC'18). doi.org/10.25344/S4QP4W

Status: GASNet-EX RMA Performance Versus MPI


- Three different MPI implementations
- Two distinct network hardware types
- On four systems the performance of GASNet-EX matches or exceeds that of MPI RMA and message-passing:
 - 8-byte Put latency 6% to 55% better
 - 8-byte Get latency 5% to 45% better
 - Better flood bandwidth efficiency, typically saturating at $\frac{1}{2}$ or $\frac{1}{4}$ the transfer size

8-Byte RMA Operation Latency (one-at-a-time)



GASNet / PAW-ATM Nov. 2019 / Paul H. Hargrove

Uni-directional Flood Bandwidth (many-at-a-time)



GASNet-EX results from v2018.9.0 and v2019.6.0. MPI results from Intel MPI Benchmarks v2018.1. For more details see Languages and Compilers for Parallel Computing (LCPC'18).

doi.org/10.25344/S4QP4W

More recent results on Summit here replace the paper's results from the older Summitdev.



Active Messages (AM)

- AM is a *restricted* form of remote procedure call
 - Executes code (handler) on a remote node
 - Request handler may *only* send back an optional Reply
 - No other communication is permitted in AM handlers
- Request and Reply APIs take
 - Integer “handler index” (which table entry to run)
 - Zero or more 32-bit integer arguments
 - Optional bulk data payload
- These arguments and payload are provided to handler

AM Payloads / Problem Statement

- Three “categories” depending on presence and handling of the optional payload
 - Short: No payload
 - Medium: Payload buffered by implementation
 - Long: Payload delivered to client-specified address
AM delivery coupled with RMA payload Put
- AM Long presents implementation challenge: to both...
 - Leverage RDMA h/w for most efficient payload transfer
 - Ensure the payload is in-place before handler runs

AM Long Protocol Tuning


Evaluating alternatives on the Cray XC's Aries network

- Latency between initiator and target
 - Overheads (CPU use on initiator and target)
 - Bandwidth between initiator and target
- } “L”, “o” and “g” of LogP model
- Sensitivity to attentiveness
 - Is more than one library entry needed to complete?
 - Timely signaling of local completion
 - Allows initiator to reuse or free payload source memory

AM Long Protocols 1

Ordered Networks (including selectively ordered)


- Simple if available, but seldom “free”
- Aries provides only at cost of defeating multi-pathing



AM Long Protocols 2

Initiator Chaining


- Simplest but also poor by most of our metrics
- Synchronous variant (aka “put-sync-send”) ties up injector until Put is complete [aries used this previously]
- Asynchronous variant relies on attentiveness



AM Long Protocols 3

Rendezvous Get


- Initiator adds source address to the message envelope
- Target uses an RDMA Get for the payload
- Adds a round-trip latency
- Also delays notification of local completion



AM Long Protocols 4


Target-Side Reassembly

- Subject of this presentation
- At very high level:
 - Payload and envelope injected into *unordered* network
 - Logic at target is *tolerant* of any reordering



Target-side Reassembly

- Both envelope and payload sent without delay
 - Can leverage multiple paths of networks like Aries
 - No attentiveness problem on either end
 - No network round-trips, and thus no stalls
 - No delays in signaling of local completion
 - However, additional network-specific metadata is required to allow target to match them



Target-side Reassembly On Aries

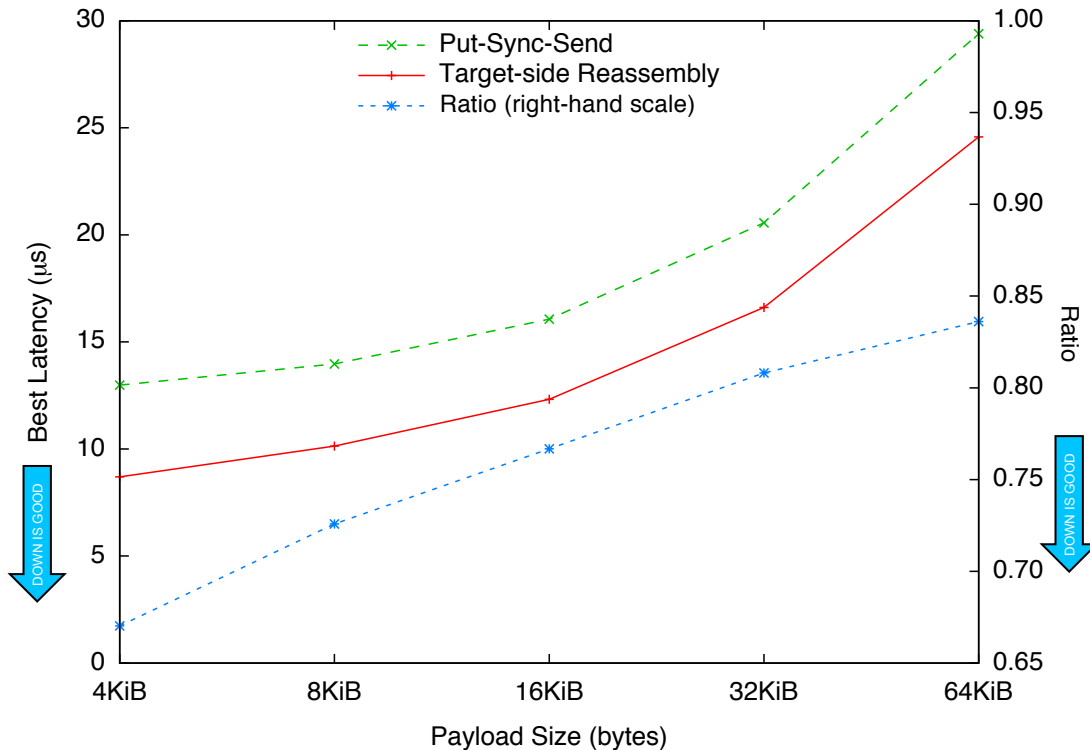
- Target needs a “nonce” to match AM envelope + payload
- Fairly simple to add a field to AM envelope
 - But we didn’t actually need to in this case
 - An existing buffer management field “fits the bill”
- Not always simple to deliver a nonce with an RDMA Put
 - Other network APIs have “Put with immediate data”
 - uGNI API for Aries has a 32-bit source identifier
 - Under software control, allowing us to steal some bits

Performance Results

- Measurements on NERSC's Cori Phase II
 - Cray XC30
 - 1.4GHz Xeon Phi CPUs
- First results: AM Long ping-pong latency
 - At most one message in flight at any time:
 - . Node 0 sends Request of given size
 - . Node 1 issues Reply of same size
 - Report average time to complete many iterations

AM Long Ping-Pong Latency

- Figure shows the range 4KiB – 64KiB
- At 4KiB
 - Old: 13.0 μ s
 - New: 8.7 μ s
 - Ratio: 33% faster
- 4KiB – 1MiB range: reductions between 3.7 μ s and 6.1 μ s




Measuring One-way Injection Overhead

- Where put-sync-send must block for payload RDMA, target-side reassembly can return immediately
 - Reduced injection overhead → more overlap
- Report the average time for many repetitions of

```
start timer
for (int i=0; i < d; ++i)
    gex_AM_RequestLong(..., size, ...);
end timer
drain network
```
- Report for various values of queue depth d and $size$

Reduced Injection Overhead

- Time to inject d AMs of a given *size*
 - d : on the x-axis
 - size*: color & symbol
 - old/new: dashed/solid
- Average injection time is up to **50x** faster
- Remains faster when flow-control sets in



Conclusions

- Presented several algorithms for AM Long on Aries
- Identified “target-side reassembly” as most promising
- Overcame challenge of coupling nonce with payload
- Presented microbenchmarks showing improvement
- We believe this algorithm is a good choice for other networks with similar properties
- New implementation released in GASNet-EX 2019.9.0

THANK YOU

gasnet.lbl.gov

doi.org/10.25344/S4PC7M

Acknowledgements

- This research was funded in part by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration.
- This research used resources of the National Energy Research Scientific Computing Center, a DOE Office of Science User Facility supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC02-05CH11231.

BACKUP SLIDES

AM Long Ping-Pong Bandwidth

- New algorithm is shown in **RED**
- Previous algorithm is shown in **GREEN**
- **BLUE** is their ratio
 - As large as 1.49
- A different algorithm is used below 4KiB

