


# Exploring Hybrid MPI+Kokkos Tasks Programming Model

Samuel Khuviz<sup>1</sup>, Karen Tomko<sup>1</sup>, Jahanzeb Hashmi<sup>2</sup>, Dhabaleswar K. Panda<sup>2</sup>


<sup>1</sup>Ohio Supercomputer Center

<sup>2</sup>Department of Computer Science and Engineering, The Ohio State  
University




# Outline

- *Introduction*
- *GEMM*
  - *Implementation*
  - *Results*
- *Graph500*
  - *Implementation*
  - *Results*
- *Limitations*
- *Summary*




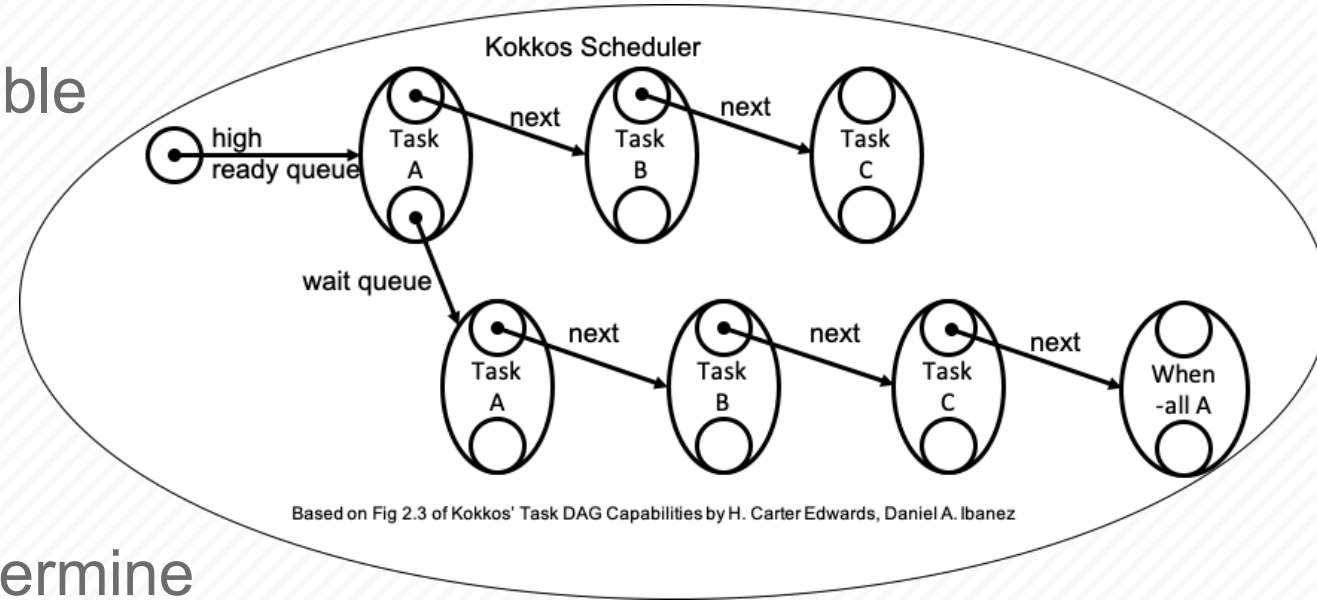
# Motivation

- Most current parallel implementations use data parallelism (MPI/MPI+OpenMP).
- Recent studies show that these models may not be well suited for emerging application.
- Task-based parallelism has potential to perform better at large-scale.
- Well-implemented task-based code will provide additional work to resources that would be otherwise underutilized.
- Most legacy scientific codes use MPI, so viable approach is to port sections of code to new models (i.e., PGAS or tasks).





# Background

- Kokkos is a C++ library for writing portable applications across different computing architectures.
- Uses a shared memory model.
- Allows for a task-based approach to parallelization.
- Tasks enqueued to a scheduler.
- Future is generated by scheduler to determine when a task has completed.
- Futures can also be used to define dependences between tasks
- A when-all entity is a future that completes when multiple tasks have completed.




# Kokkos GEMM Implementation

- Uses tasks for shared memory parallelism.
- Main function spawns 1 task with a triply nested for loop.
- Each iteration of loop launches one task at a time on each subblock of the matrices to compute the product on that subblock.


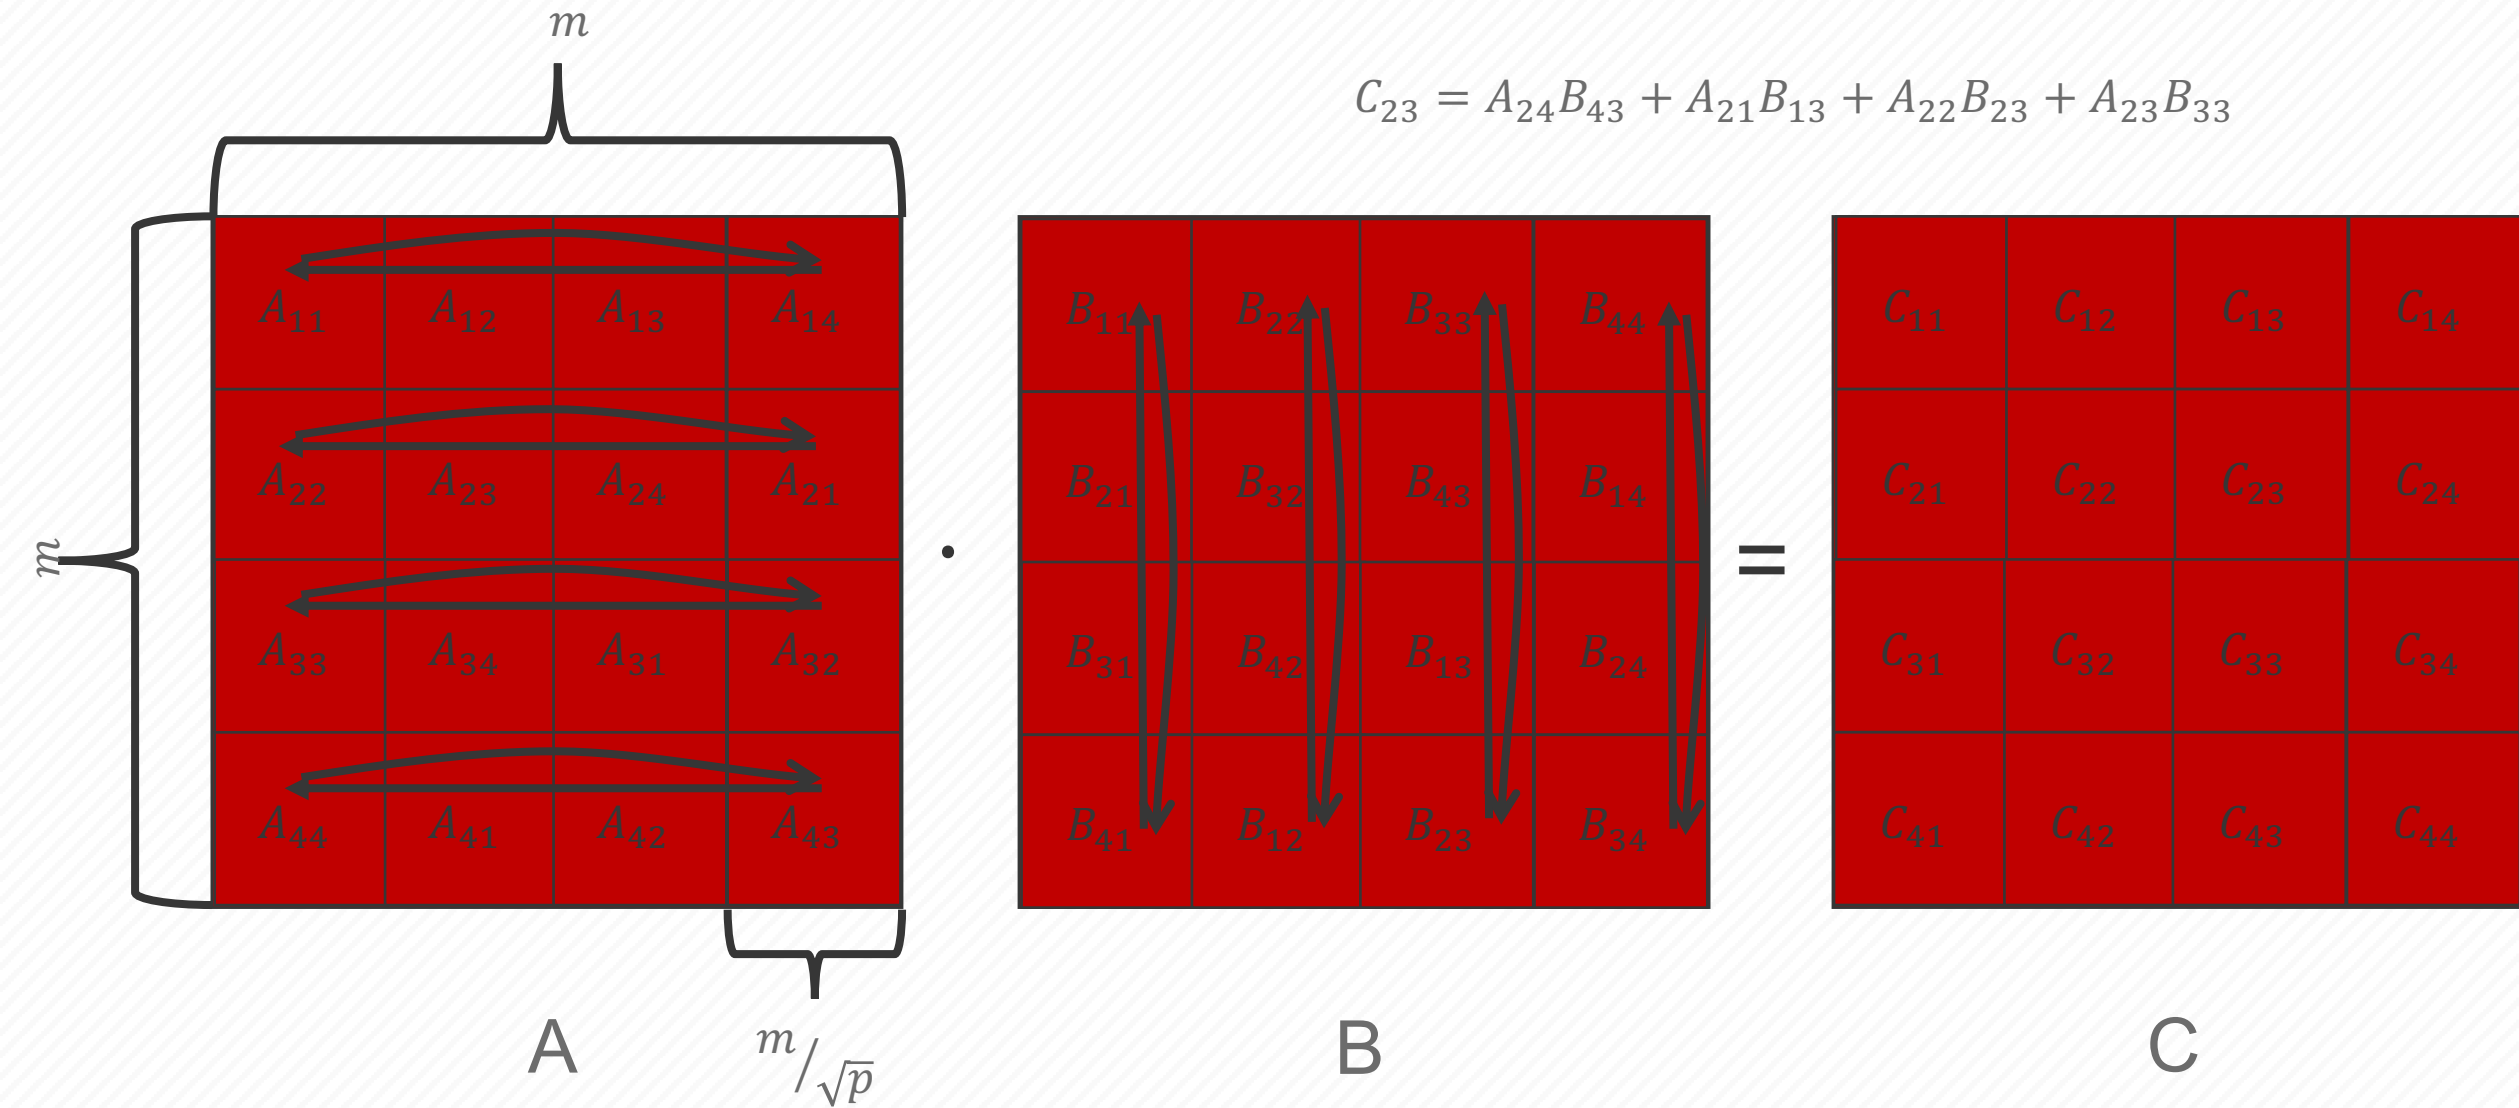


# MPI+Kokkos GEMM Implementation

- Add MPI in order to scale on multiple nodes.
- Use Cannon's algorithm since it scaled well and keeps memory usage constant as the number of ranks grows.
- No modifications to tasks since each GEMM task performs computation on local subblocks.




# Cannon's Algorithm on $p = 16$ processes



# GEMM Code on 1 node

Kokkos Tasks	Runtime (in sec)	Speedup	Efficiency
1	1223	1	1
2	681	1.80	0.90
4	336	3.64	0.91
8	174	7.02	0.88
16	97	12.61	0.79
32	58	21.09	0.66
40	50	24.46	0.61

- Problem size of 20,000 and block size of 256.
- Run on 1 rank.
- Parallel efficiency is below ideal efficiency.






# GEMM Code over multiple nodes


Nodes	Runtime (in sec)	Speedup	Efficiency
9	72.5	1	1
16	43.5	1.67	0.94
25	27.7	2.62	0.94
36	20.4	3.55	0.89
64	11.8	6.15	0.87

- Problem size of 45,000 and block size of 256.
- Run with 1 rank/node and 40 tasks/rank.
- Adding MPI enables larger problems to be solved than Kokkos-only.
- Scales well up to 64 nodes (or 2560 tasks).

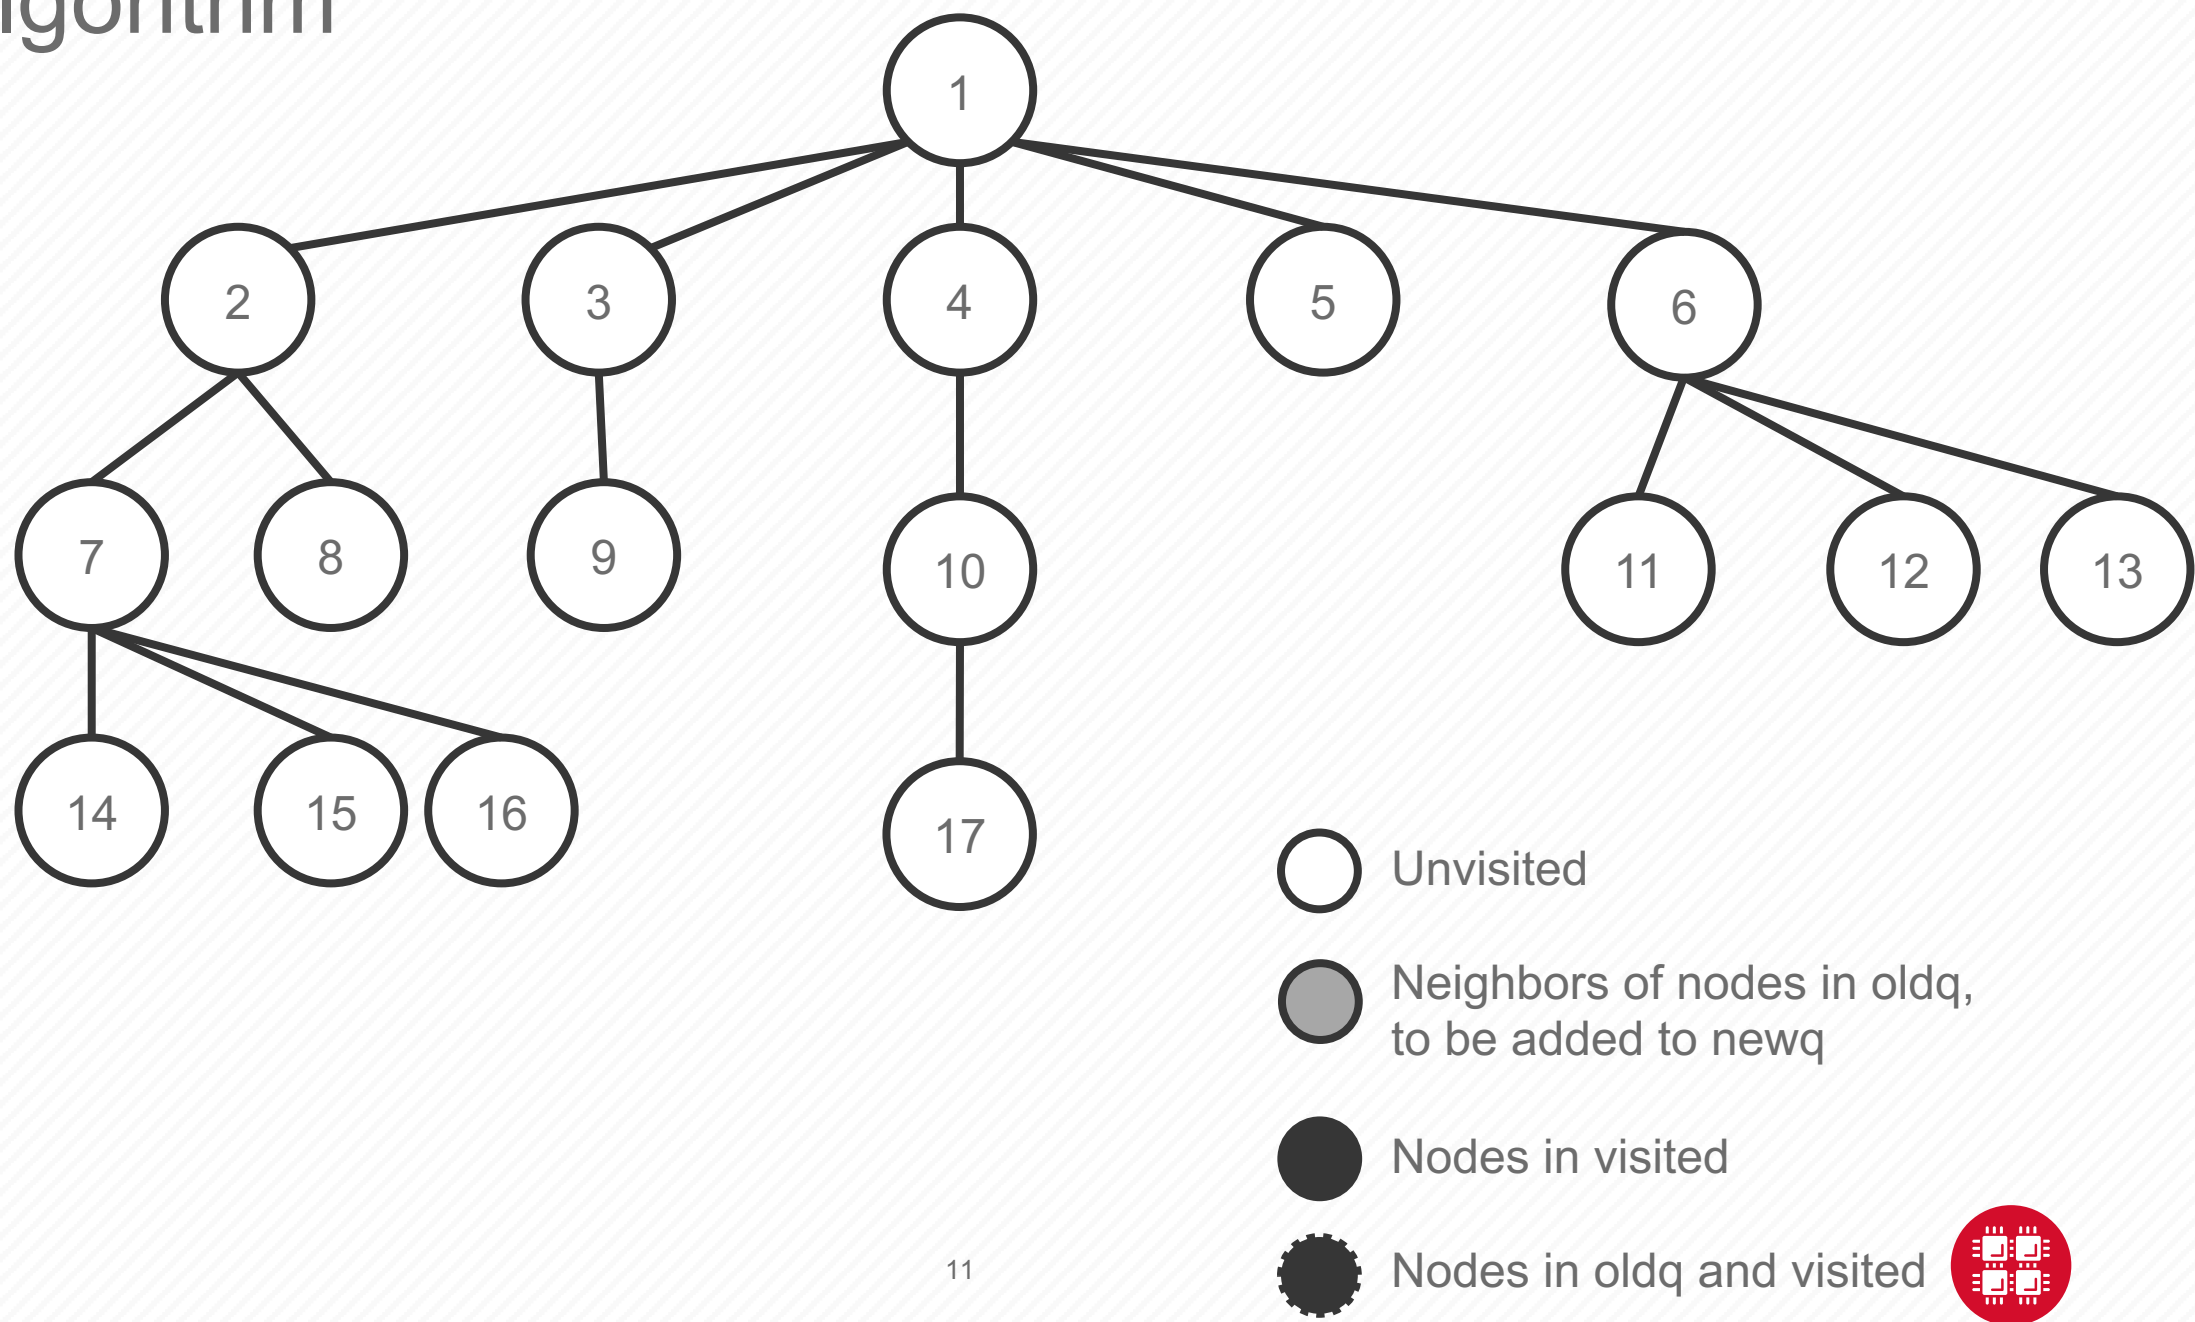


# MPI Graph500 Implementation

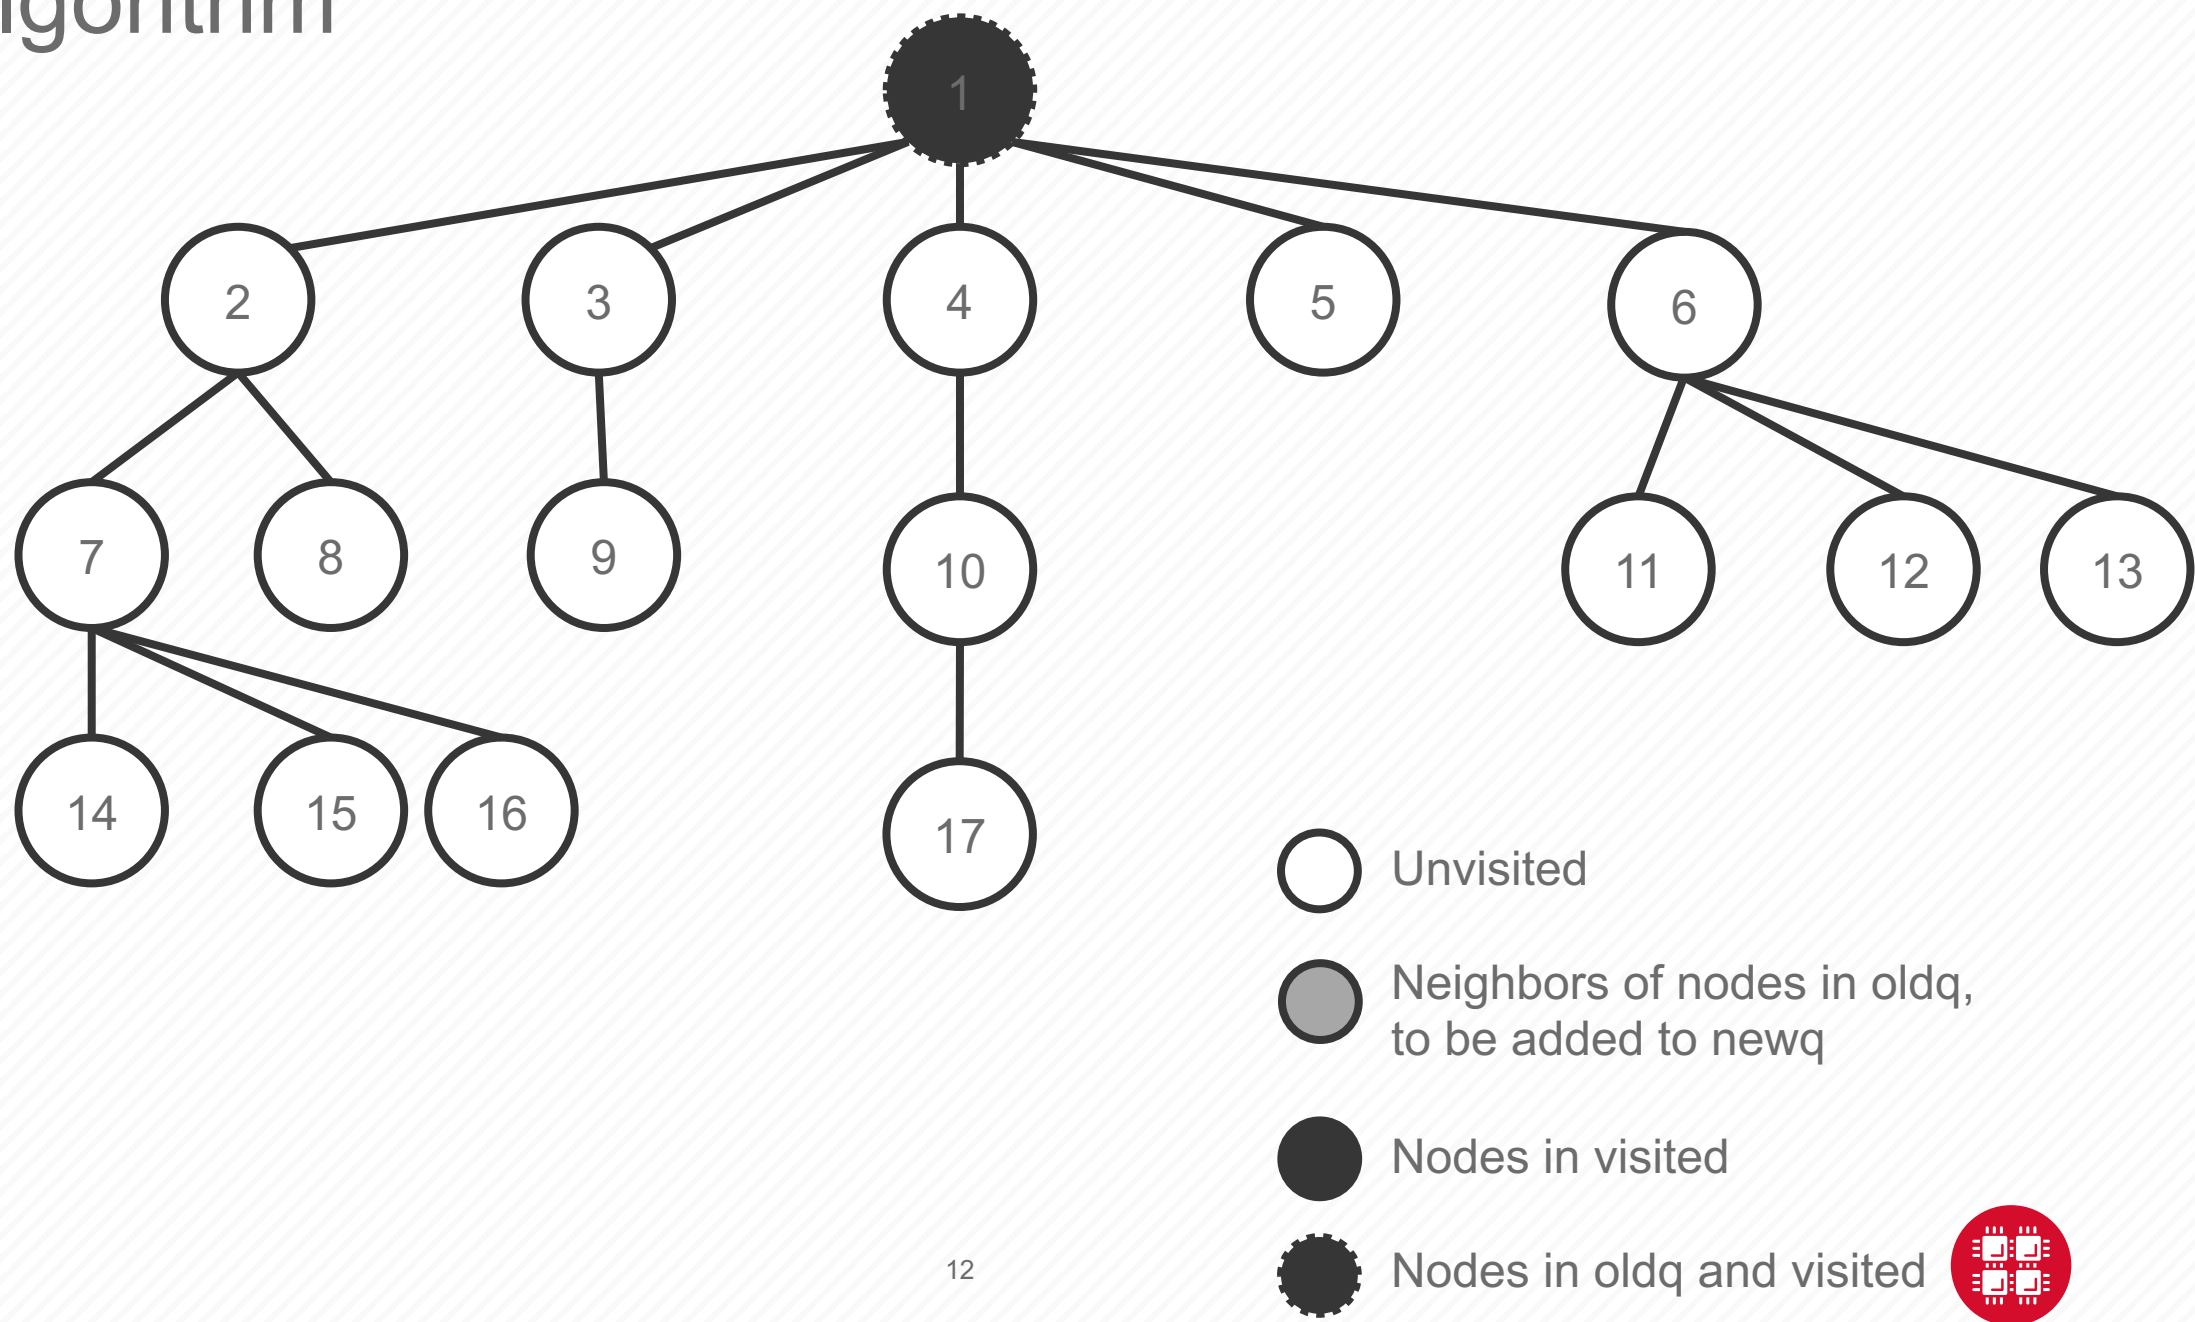
- Breadth first search (BFS) function is passed graph, root vertex, and visited array.
- Key data structures:
  - `oldq`: array of vertices visited in previous iteration,
  - `newq`: array of neighbors of vertices in `oldq`,
  - `visited`: array of vertices that have been visited,
  - `pred`: array that maps neighbor vertex to source vertex.



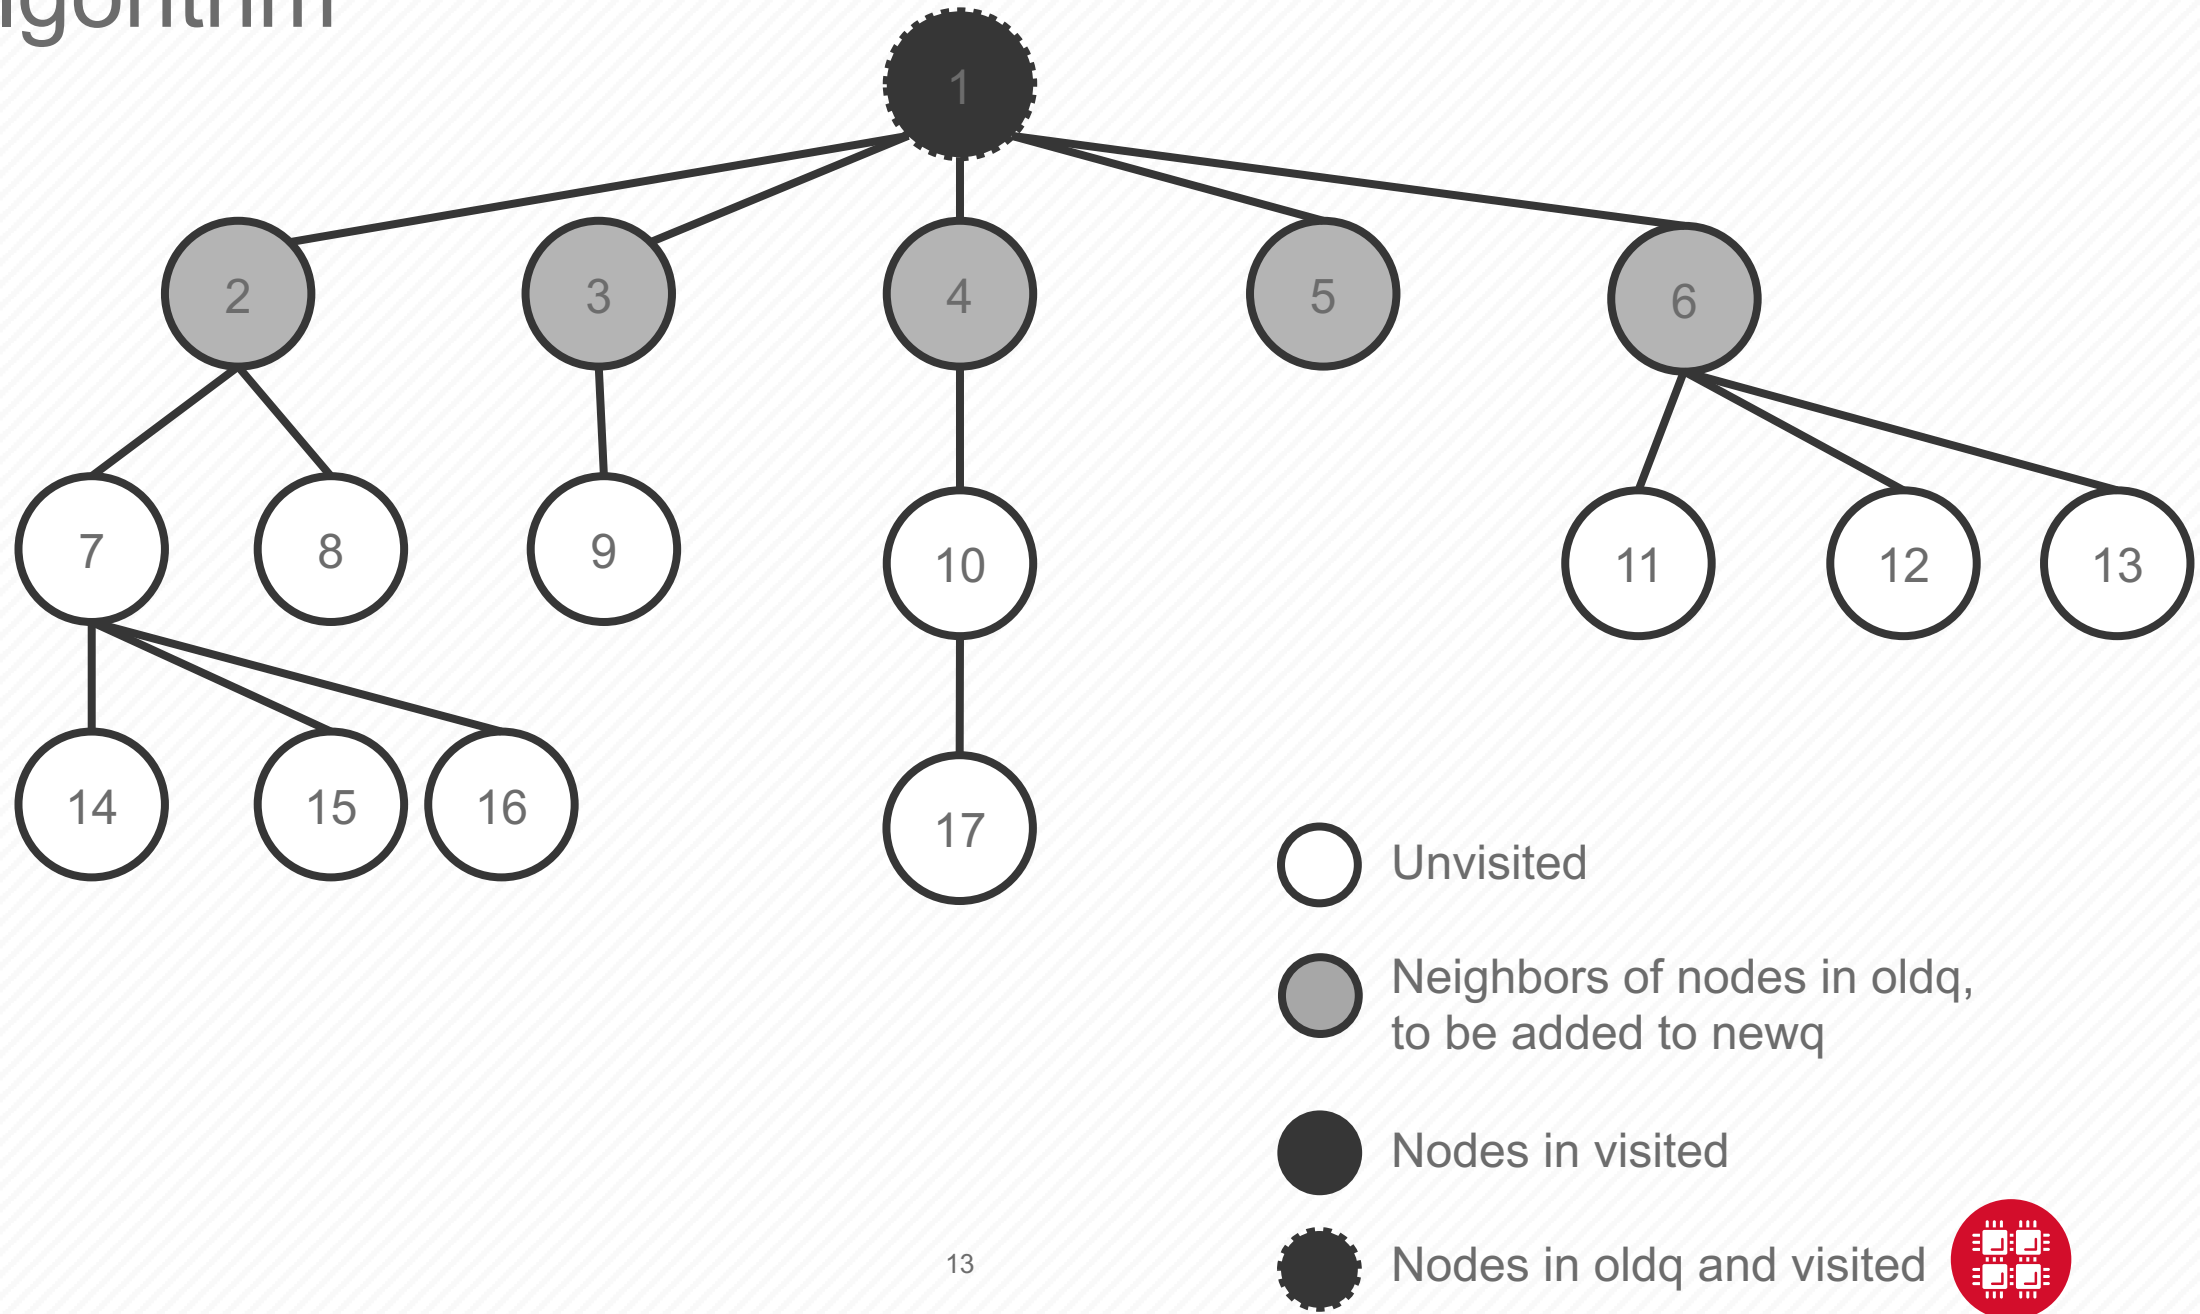
# BFS Algorithm



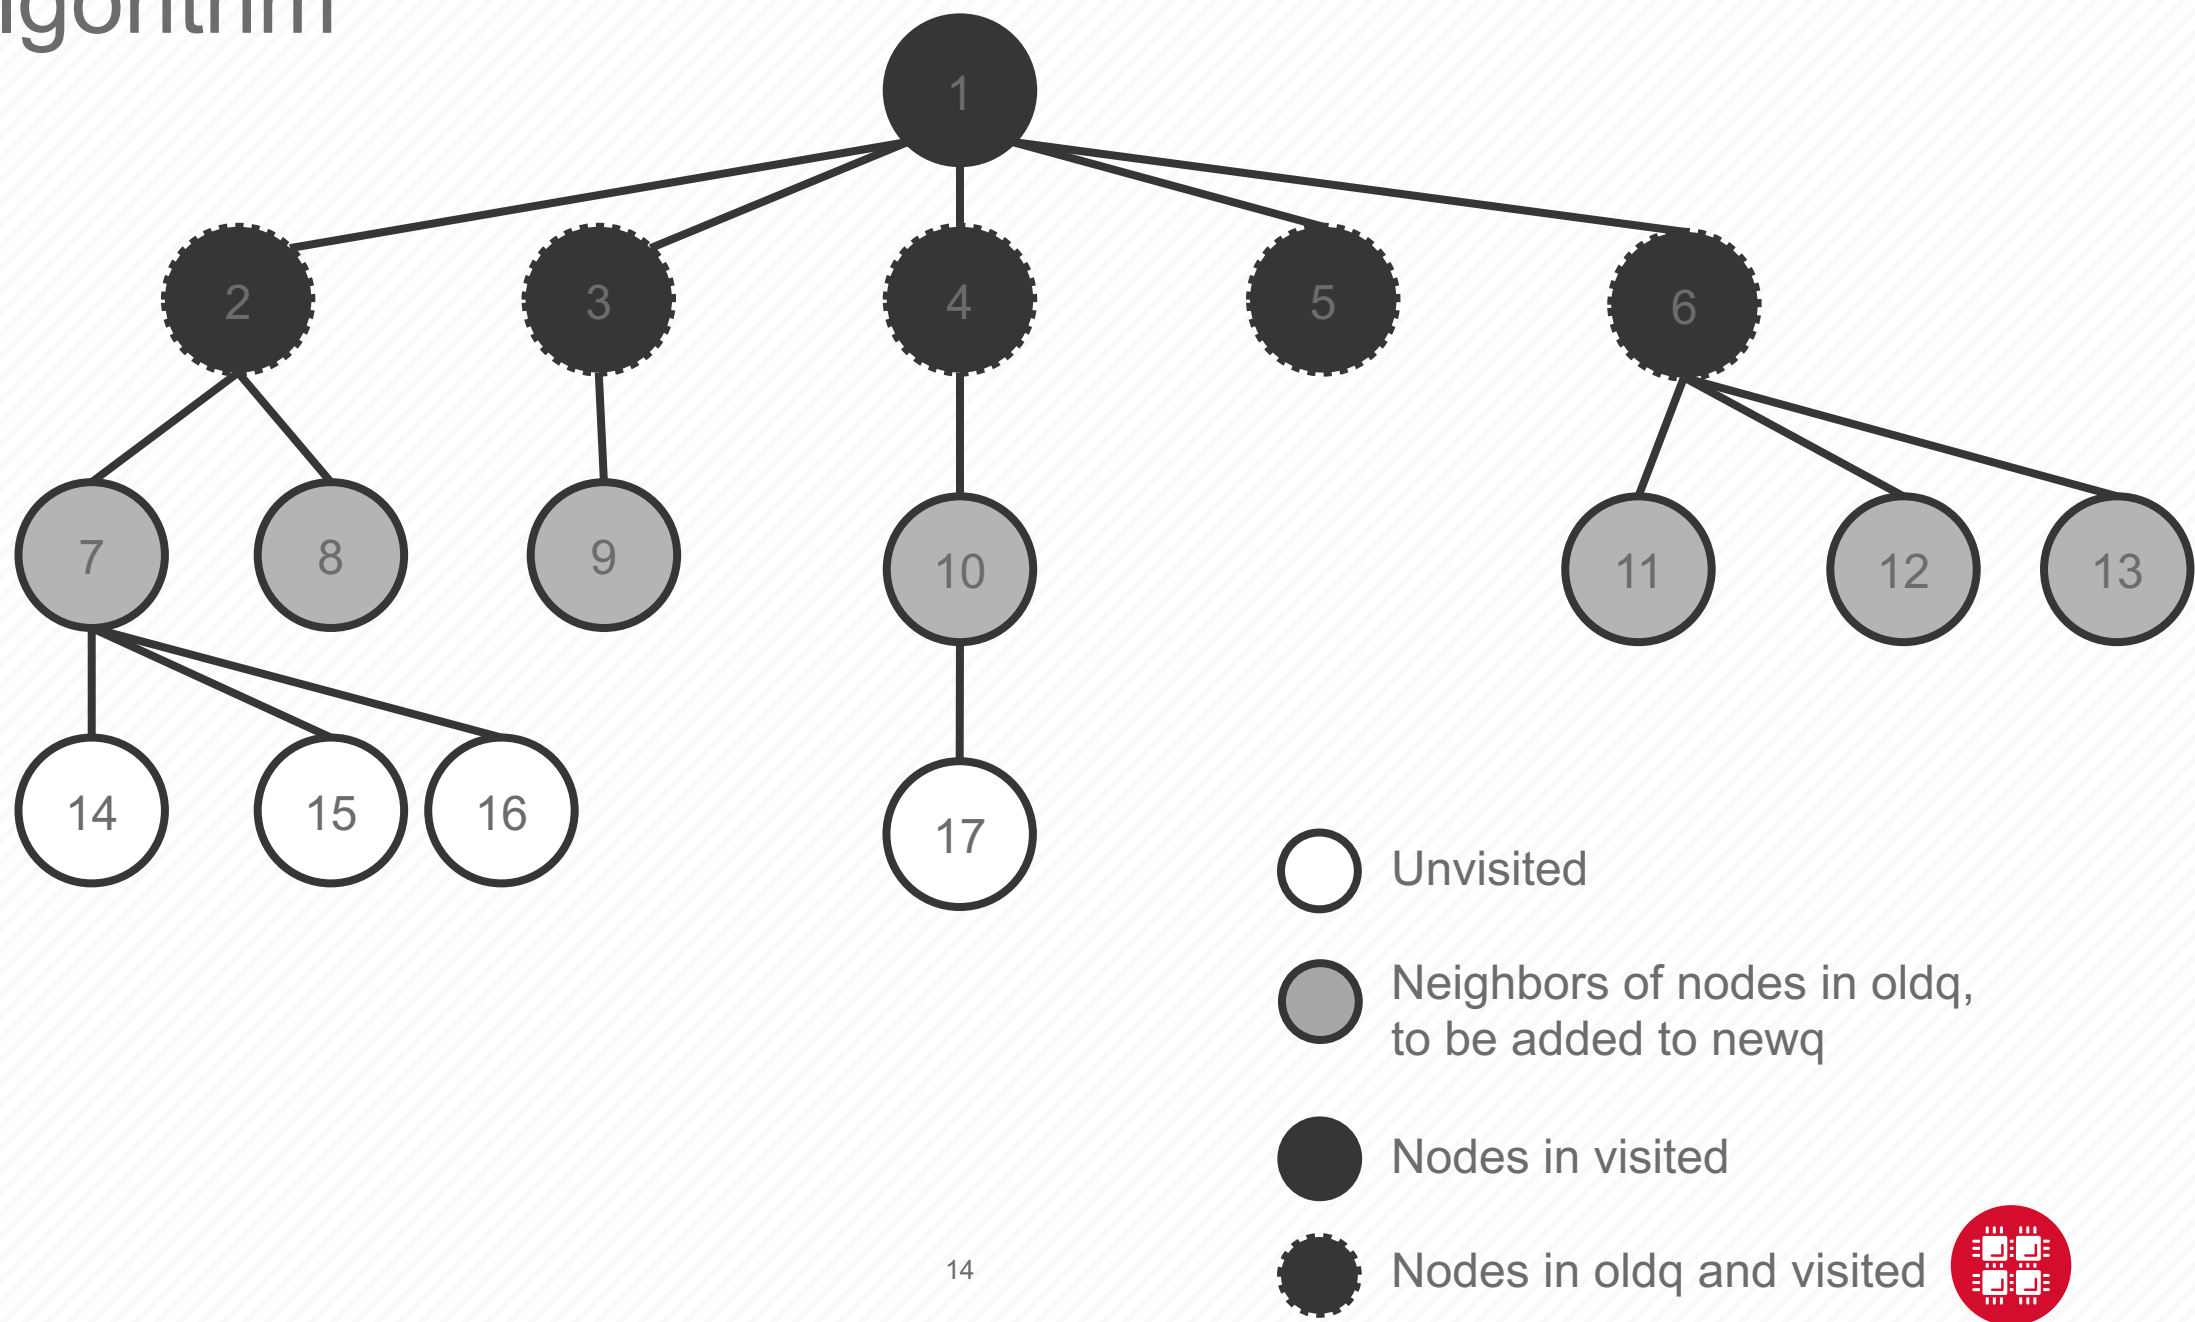
# BFS Algorithm



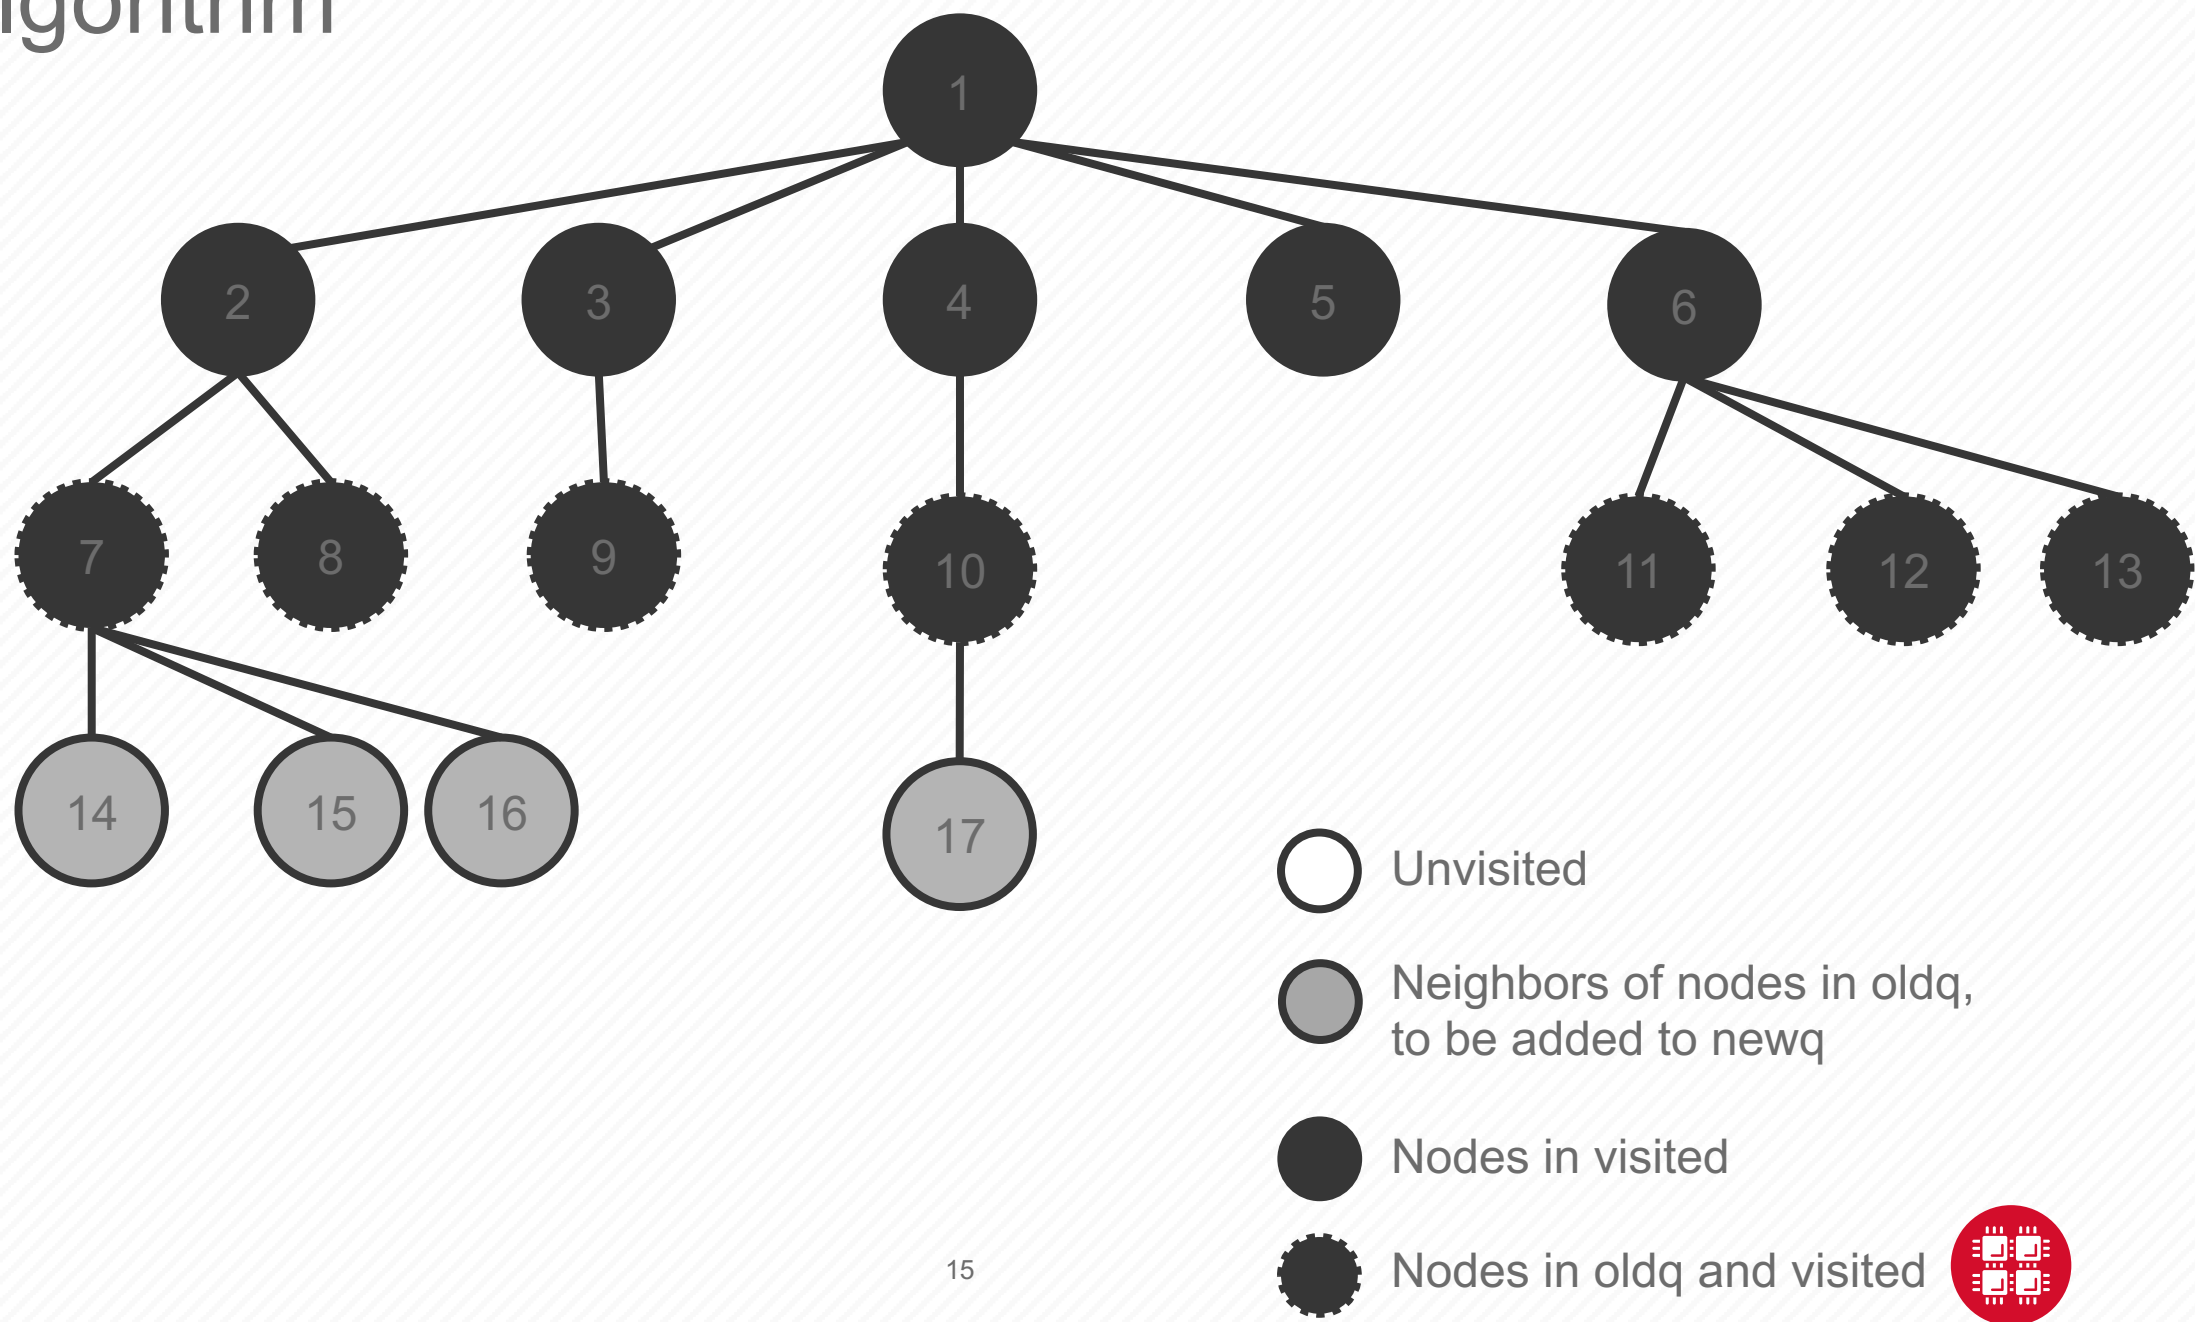
# BFS Algorithm



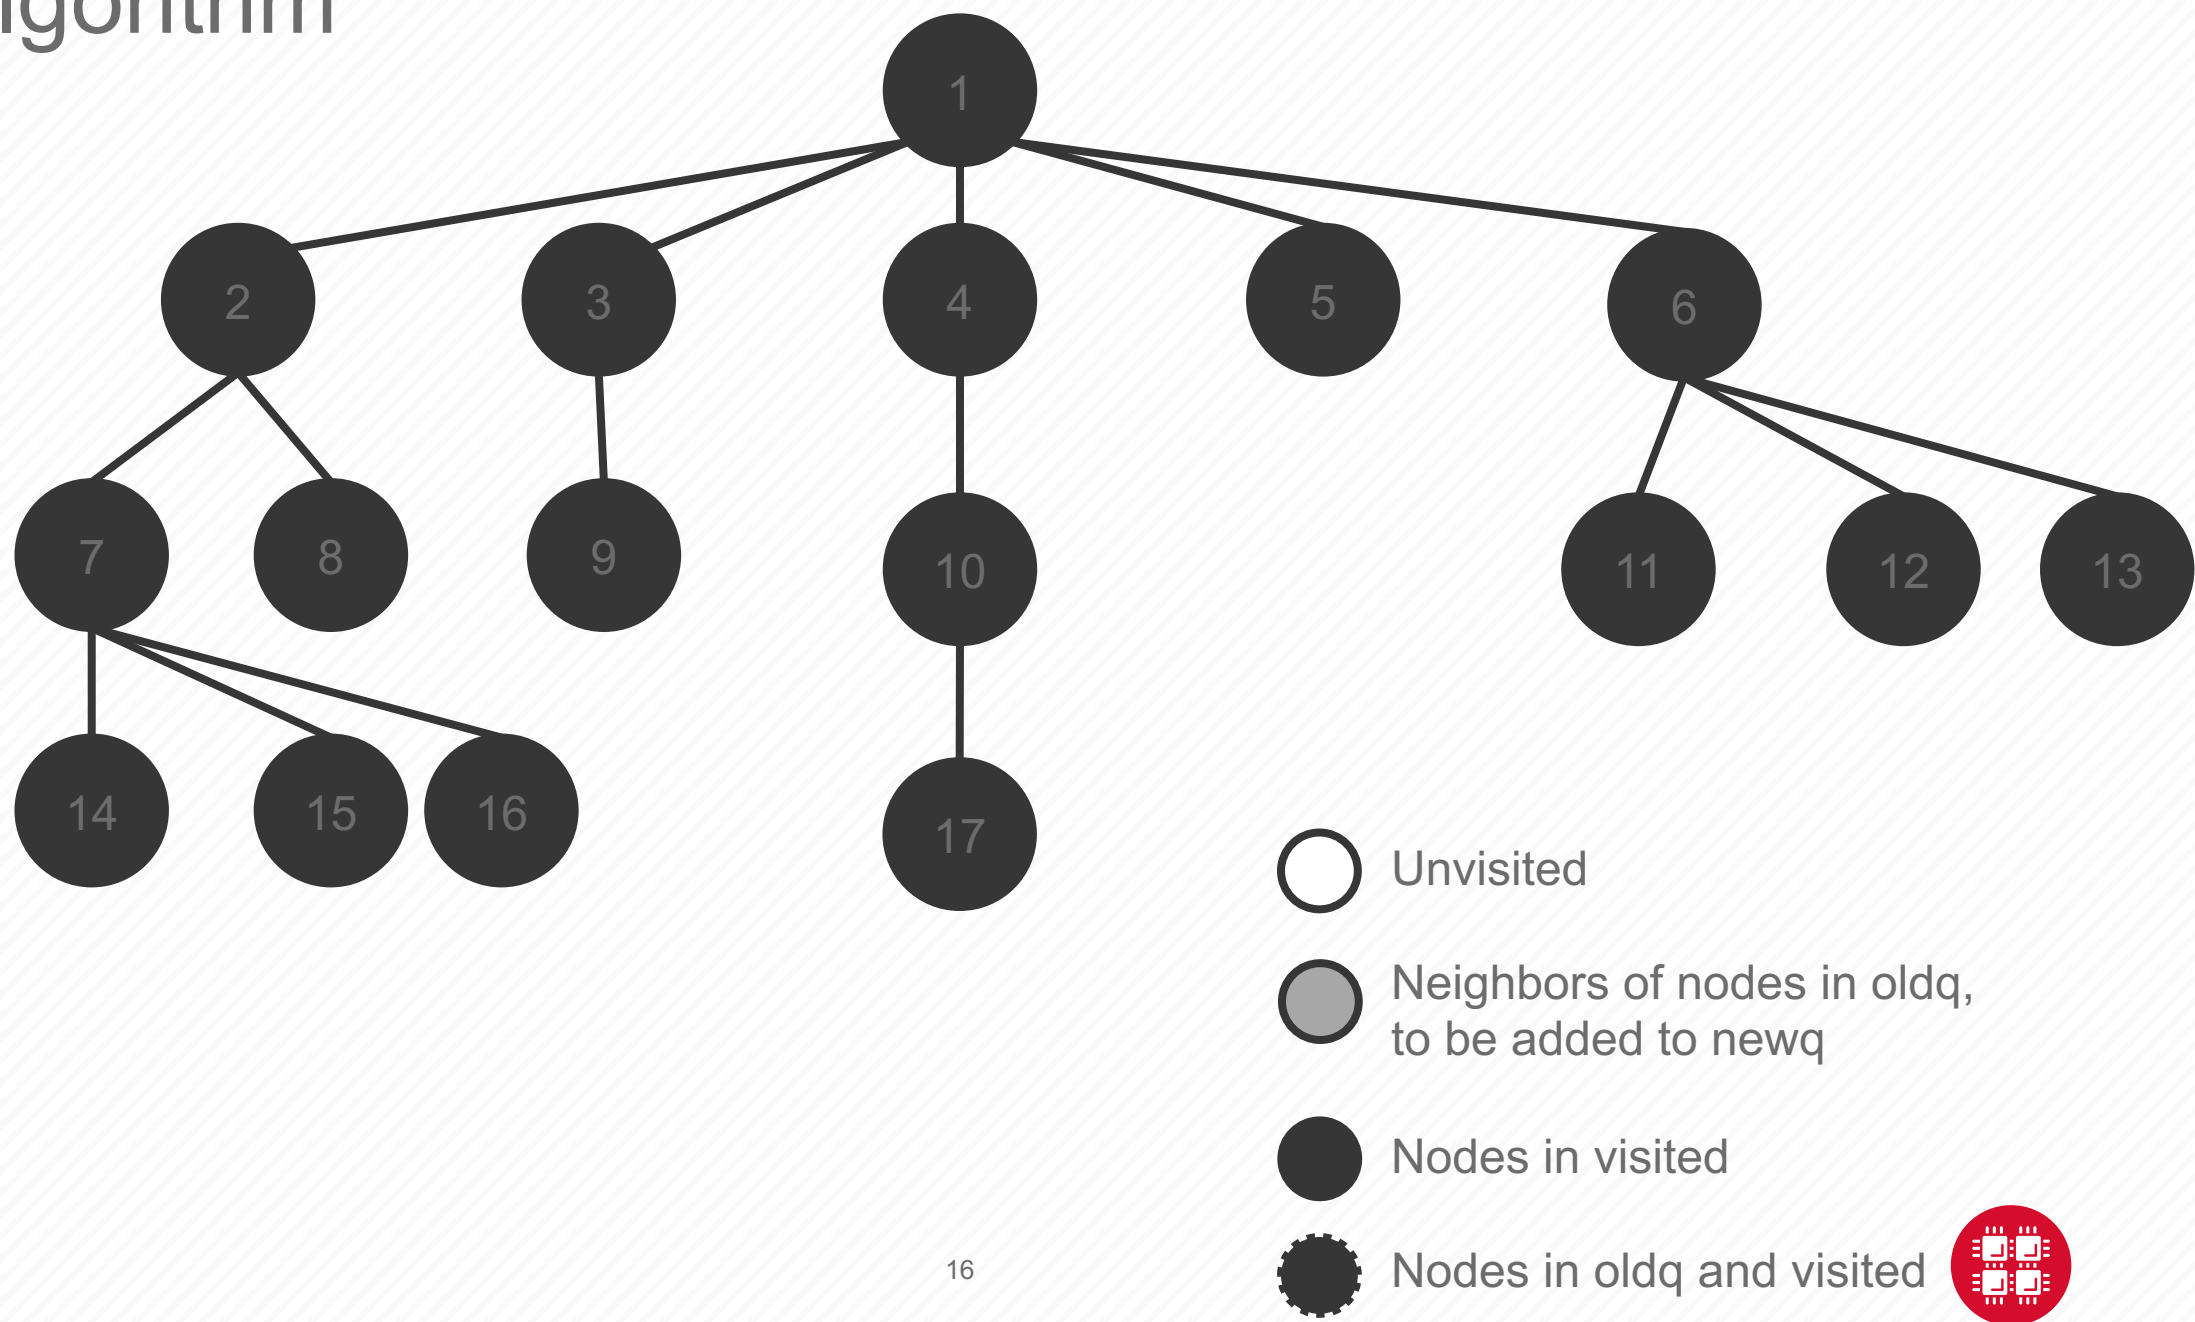
# BFS Algorithm



# BFS Algorithm




# BFS Algorithm





# MPI+Kokkos Graph500 Implementation (locking)


- Use MPI implementation as starting point.
- Each process spawns NT tasks:
  - $NT - 1$  processing tasks that iterate over a subset of `oldq`,
  - 1 receive task that waits for data from other ranks and processes it as it receives it. Runs until all processing tasks send a signal that they have completed.
- All tasks modify `visited`, `pred`, and `newq` arrays so updates are locked to avoid race conditions.
- `newq_count` atomic variable used to track elements in `newq`.



# Locking Graph500 Results


Tasks	BFS	Speedup	Efficiency	Locks
2	48.7	2.0	1.00	3.78
4	106.4	0.92	0.23	53.01
8	141.2	0.69	0.086	88.58
16	162.0	0.60	0.038	92.49
32	248.1	0.39	0.012	161.96
40	259.2	0.38	0.0094	162.85

- SCALE=20.
- Running on 1 MPI process.
- Increasing number of tasks increases runtime.
- Time spent in locks also increases with number of tasks.
- Significant percentage of slowdown caused by time spent in locks.



# MPI+Kokkos Graph500 Implementation (non-locking)


- Removes all locks and atomics.
- Each task maintains its local version of `visited`, `pred`, and `newq` arrays.
- Once all tasks have completed, local arrays from all tasks on a rank are merged into global arrays.



# Non-locking Graph500 Results

Tasks	BFS	Speedup	Efficiency
2	25.5	2.0	1.00
4	18.7	2.7	0.68
8	14.2	3.6	0.45
16	11.3	4.5	0.28
32	10.6	4.8	0.15
40	10.2	5.0	0.13


- Better absolute runtimes than locking implementation.
- Scales better than locking implementation, but still poorly.



# Non-locking Graph500 Results

Processes	Processing Tasks	Runtime (in s)
1	16	11.6
2	8	11.1
4	4	7.5
8	2	7.6
16	1	4.7


- Test different combinations processes and tasks for a total of 16 processing tasks.
- Increasing processes/ decreasing tasks improves performance.
- Best performance with 16 processes and 1 processing task.



# Imbalance Percentage for Non-locking Graph500 Results

Tasks	Max Imbalance	Mean Imbalance
4	0.857	0.471
8	0.946	0.471
16	0.956	0.472
32	0.926	0.476
40	0.956	0.572


- Percentage of time task spends waiting for other tasks to complete.
- For each iteration, compute difference in runtime between slowest and fastest tasks, then divide by time spent in slowest task.
- On average, over all iterations, ~50% of task time is spent in waiting.
- For some iterations, imbalance can be as high as 96%.



# Graph500 Results Comparison


Processes	MPI	Locking MPI+Kokkos	Non-locking MPI+Kokkos
1	10.5	254.5	10.2
2	13.4	83.2	12.3
4	7.5	34.1	7.8
8	4.3	15.0	5.8
16	2.9	7.4	4.4
32	2.3	N/A	N/A

- Non-locking performs better than locking.
- Best performance with MPI-only.



# Limitations


- Profiles of Graph500 show more time is spent in MPI communication when using Kokkos (more data is sent with a lower call rate).
- Non-locking implementation of Graph500 performs additional work since multiple tasks may visit the same node.
  - For SCALE=20, 652,549 unique vertices but 2,515,974 to 4,077,491 vertices visited across all tasks (3.86 to 6.25 times more vertices).
- Can only check if all tasks have completed with a barrier.
  - ~50% of processing task time spent waiting.
  - Wait must complete for all tasks before merging local arrays into global arrays.
  - It would be beneficial to overlap queue processing in the tasks with merging local arrays by merging as tasks complete.






# Summary

- Demonstrated feasibility of hybrid MPI+Kokkos tasks codes
- Demonstrated that hybrid MPI+Kokkos tasks model can efficiently solve larger problem sizes than Kokkos tasks alone
- Evaluated scaling of 2 example hybrid MPI+Kokkos codes
- Identified challenge to adopting this task model



# Acknowledgements

- We would like to thank the Ohio Supercomputer Center for access to the Pitzer cluster.
- This research is supported in part by NSF grants #1931537 and #1565414.





# OH·TECH

Ohio Technology Consortium  
A Division of the Ohio Department of Higher Education

 [info@osc.edu](mailto:info@osc.edu)

 [twitter.com/osc](https://twitter.com/osc)

 [facebook.com/ohiosupercomputercenter](https://facebook.com/ohiosupercomputercenter)

 [osc.edu](http://osc.edu)

 [oh-tech.org/blog](http://oh-tech.org/blog)

 [linkedin.com/company/ohio-supercomputer-center](https://linkedin.com/company/ohio-supercomputer-center)

