

Shared Memory HPC Programming: Past, Present, and Future



Bill Carlson


IDA Center for Computing Sciences

November 13, 2017

Disclaimer

This presentation reflects the personal views of the author.

*These views may or may not be held by the author's
employer or its sponsors.*



Our Problem in 1993

How do we program this? And get good performance?

AC for the Cray T3D

- ❖ An outgrowth of our work on CM5
- ❖ Shared memory on a distributed memory machine
 - ❖ “dist” keyword is the only syntax change
 - ❖ Performance high from special hardware on T3D
 - ❖ Faster than “shmem” library, due to low overhead


$$\text{UPC} = \text{AC} + \text{Split-C} + \text{PCP}$$

- ❖ Collaboration with UC Berkeley and LLNL
 - ❖ Takes “shared” from AC’s “dist”
 - ❖ “strict” and “relaxed” shared memory semantics
 - ❖ Split barriers: “notify/wait”
 - ❖ Locks
 - ❖ Adds several data distributions

PGAS: Expanding the collaboration



- ❖ SHMEM Library
- ❖ CoArray Fortran
- ❖ Global Arrays
- ❖ Titanium



DARPA HPCS Program

3 New PGAS Languages!

- ❖ Fortress: Implicit Parallelism, Strong Types
 - ❖ Cool look: like math in both ASCII and Unicode
 - ❖ Effort ended in 2012
- ❖ X10: Java-like syntax, asynchrony, locales
 - ❖ Going strong but, not much “HPC”
- ❖ Chapel: Separate Parallelism and Locality
 - ❖ Annual Conference, open source distribution

Post-HPCS PGAS

- ❖ Habanaro C and UPC++ (Rice)
- ❖ UPC++ (Berkeley)
- ❖ CoArray C++ (Cray, EPCC)
- ❖ HPX (C++/11,14, LSU, FAU)
- ❖ XcalableMP (Tskuba)
- ❖ GASPI (Fraunhofer)
- ❖ OpenSHMEM
- ❖ More every time one looks


HPC at a Crossroads

- ❖ Path to Exascale is underway:
 - ❖ System complexity increasing
 - ❖ Huge increase in OPs
 - ❖ Lesser increases in memory and communication performance
- ❖ Application development is getting harder, not easier



PGAS at a Crossroads

- ❖ Many implementations exist of PGAS techniques, essentially all platforms
- ❖ Provide a wealth of programming metaphors
- ❖ Performance has been shown to be very good
 - ❖ A number of cases which exceed best message passing code
 - ❖ Because you have a wider choice of algorithms and synchronization
- ❖ Programmer “base” is
 - ❖ (somewhat) small, and
 - ❖ (somewhat) static



Our Problem in 2017

How do we program this? And get good performance?

Key problem: Random Communication

- ❖ Many PGAS apps access shared memory randomly
 - ❖ Problematic on modern systems, due to relatively high cost of message initiation
- ❖ We have an approach called “exstack/Conveyors”:
 - ❖ Generate many more accesses than threads
 - ❖ Sort them locally, send in batches
 - ❖ Take received messages, process locally
 - ❖ Repeat until done
 - ❖ Several synchronization styles supported

exstack/Conveyors Analysis

- ❖ Performance surprisingly good for most apps
 - ❖ Most have plenty of parallelism in their accesses
 - ❖ Most tolerate the significant added latency
- ❖ Programability is a complete mess
 - ❖ Ugly, lengthy, complex looping structures in code
 - ❖ Have to packetize work by hand
- ❖ We want better, much like our previous “AC” work
 - ❖ We have tried to improve this and not found key
 - ❖ Aware of related efforts, we don’t see key there either
 - ❖ Yet!

Thought Questions for Today

- ❖ How important is HPC performance?
 - ❖ And what does “performance” mean?
- ❖ How important is HPC programability?
 - ❖ And what does “programability” mean?
- ❖ How can we deal with HPC system complexity?
- ❖ Should we expand the PGAS horizons beyond HPC?

HPC Performance

- ❖ Performance has always been *critical* to HPC
 - ❖ By definition. If not, why go to the trouble. What matters:
 - ❖ Application results per unit time (e.g., day)
 - ❖ Application results per unit cost (e.g., \$, Watt)
- ❖ Ways to improve application performance
 - ❖ Understand performance limiters, then
 - ❖ Write new code which gets around limiters
 - ❖ Includes both tweaks and new algorithms
 - ❖ Use new systems which address limiters
 - ❖ Hardware, Compilers, OS, etc, etc.

HPC Programability

- ❖ Programability has always been *debated* in HPC
 - ❖ By definition: Performance is *critical*. What matters:
 - ❖ Application performance achieved per unit cost
- ❖ Ways to improve programability
 - ❖ Apply more or better programmers
 - ❖ Apply more or better tools for programmers
 - ❖ Apply different programming techniques
- ❖ Debate is always about mix and extent of efforts

“An unwritten program has ZERO performance”

—Me

Key HPC Complexity

Multi-Level Parallelism

- ❖ Hardware is becoming increasingly hierarchical
 - ❖ Start with SMP “nodes” in distributed machines
 - ❖ Add threads within cores within processors
 - ❖ GPUs and other accelerators only add to the complexity
- ❖ Two distinct issues:
 - ❖ What is shared among threads on a “node”? But not globally?
 - ❖ What controls the parallel activity on a node?

Multi-Level Parallelism?

- ❖ Some programming models urge multi-level
 - ❖ SHMEM + pthreads or OpenMP
 - ❖ Programmers then write two levels of control flow, one for across nodes, one for on nodes
- ❖ UPC supports only local and shared
 - ❖ PGAS thread per hardware thread seems about right
 - ❖ An extension was made to allow shared allocation on node

Is “HPC” the only PGAS “market”?

- ❖ Mostly yes
 - ❖ Pointless to “partition” a tiny system
- ❖ But maybe not!
 - ❖ No widely-useful model for programing SMP processors
 - ❖ Most restricted to concurrency (e.g., go)
 - ❖ PGAS could provide a path to scalable apps
 - ❖ PGAS can be powerful metaphor in programmer education

PGAS Future?

- ❖ Stay the Course?
- ❖ Another Unification?
- ❖ New Approach?



Path Forward One: Keep Pressing

- ❖ Our current languages and libraries are good!
- ❖ Our current programmers are good!
- ❖ We are growing friends all the time
- ❖ To Do List:
 - ❖ Implement github-scale sharing of PGAS utilities
 - ❖ Start work on new application areas
 - ❖ Develop curriculum

Path Forward Two: New Unification

- ❖ UPC took three smaller, locally used languages
 - ❖ And made something better than sum of parts
- ❖ Several areas to consider
 - ❖ Many C++ based PGAS efforts are underway
 - ❖ PGAS Multi-treading and asynchrony
 - ❖ Including efforts that don't know they are PGAS, yet :)
- ❖ But gaining branding and adoption is always hard


Path Forward Three: New Approach

- ❖ What's all this about Python (and friends) then?
 - ❖ No need to change the language to change programming models
 - ❖ Abundant local computation means we can (maybe) afford the overhead
 - ❖ And maybe call “native” code when needed
 - ❖ Can us any good communication system
- ❖ Why is there no “PGAS” for Python?
 - ❖ There is (or was), google “Python PGAS”
 - ❖ It was a good thought but never got supported

Consider: JavaScript and Node.js


- ❖ Node.js is a JavaScript engine based on Chrome's V8. All I/O is event driven and non-blocking
- ❖ Performance: See "Benchmarksgame"
 - ❖ JavaScript is about 10x Python, 1x go, java, 0.1x C
- ❖ Interesting model
 - ❖ Every piece of JavaScript is sequential
 - ❖ But much parallelism based on event driven "I/O"
 - ❖ Which provides good mechanisms to deal with parallelism
- ❖ My attempt is to "PGAS" this where performance penalty is acceptable

PGAS Node.js



 UPC/PGAS Process: started by SLURM

❖ Fixed program for inter-node accesses and synchronization

 Node.js Process: `fork () / exec ()` by UPC/PGAS

❖ User program, use `mmap ()` communication with 

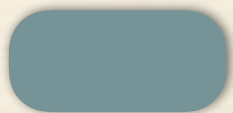
PGAS Shared Objects

- ❖ All shared data derived from a root shared object
 - ❖ Data stored in PGAS program
 - ❖ Methods are something like `get()`, `put()`, `action()`
 - ❖ Access to shared mediated by PGAS program
 - ❖ Completion through callbacks, “promises”
 - ❖ May provide the programability and performance

Histogram in JavaScript

```
let num_per = 2457;
let span = par.num_proc*num_per;

Histogram(par, num_per, function(h) {
  for (let i = 0; i < 12345*par.num_proc; i++) {
    h.update(Math.random()*span);
  }).then(h => print h.histo[0]);
```

- ❖ **par** is an object which holds parallel context, from 
- ❖ **Histogram** is a specialization of **ParallelUpdater**
 - ❖ `update()` parallel communication
 - ❖ Callback parallelizes allocation
 - ❖ Promise “`.then()`” handles completion
- ❖ Really just starting here, but I’m excited!

Closing Thoughts

- ❖ PGAS has developed a significant history
- ❖ Currently at a crossroads:
 - ❖ Concepts have proved useful in HPC
 - ❖ Struggling with current architectures
 - ❖ Adapt new approaches
- ❖ All future paths are interesting
 - ❖ But JavaScript is fun for me now, but happy to help with others!



Future Vistas for PGAS

The fun has only begun

Image Credits

- ❖ T3D Image: CC-BY-SA-2.0-fr Rama, Wikimedia: CRAY-T3D IMG 8981-82-87-89.CR2.jpg
- ❖ Expanding Image: Public Domain, "CMB Timeline300 no WMAP" by NASA Wikimedia:CMB_Timeline300_no_WMAP.jpg
- ❖ Crossroads Image: CC-BY-SA Umberto Nicoletti, flickr.com/photos/unicoletti/2851575552
- ❖ Titan Image: Public Domain, James086, Wikimedia: Titan render.png
- ❖ Final Image: CC-BY-2.0 Nicholas A. Tonelli, [flickr.com/photos/nicholas_t/6697202819](https://www.flickr.com/photos/nicholas_t/6697202819)

Backup Slides

DARPA's HPCS Program

- ❖ High ***Productivity*** Computing Systems
- ❖ Productivity: Output per unit of Input
 - ❖ Output is problems solved
 - ❖ Input is money, energy, people time
- ❖ Goal: Increase productivity of HPC by 10x:
 - ❖ Systems performance 10x for many metrics
 - ❖ Algorithm and Software developers 10x effective in making good code
 - ❖ System operators spend 1/10 effort to manage system

New Uses for HPC/PGAS

- ❖ A lot of emphasis on “Big Data”
 - ❖ How about an awesomely fast PGAS key-value store
- ❖ Machine “Deep” “learning”
 - ❖ Can PGAS allow real advances in this field
- ❖ Previously “Abandoned” HPC applications
 - ❖ Industrial uses in manufacturing
- ❖ My assertions:
 - ❖ PGAS approaches could help add new application areas
 - ❖ These are not using HPC (much) because it is too hard to program

Complexity of Next Gen HPC?

- ❖ Strong forces for higher complexity
 - ❖ Need to control energy leads to specialization
 - ❖ Accelerators like GPUs
 - ❖ Small, specialized memories
 - ❖ Communication at a distance is always limited by cost
 - ❖ ExaScale goals are pushing for large performance gains
- ❖ Some trends to lower complexity
 - ❖ Many applications can fit “on a node” or small segment of system
 - ❖ Communication bound algorithms might ignore complex parts of system